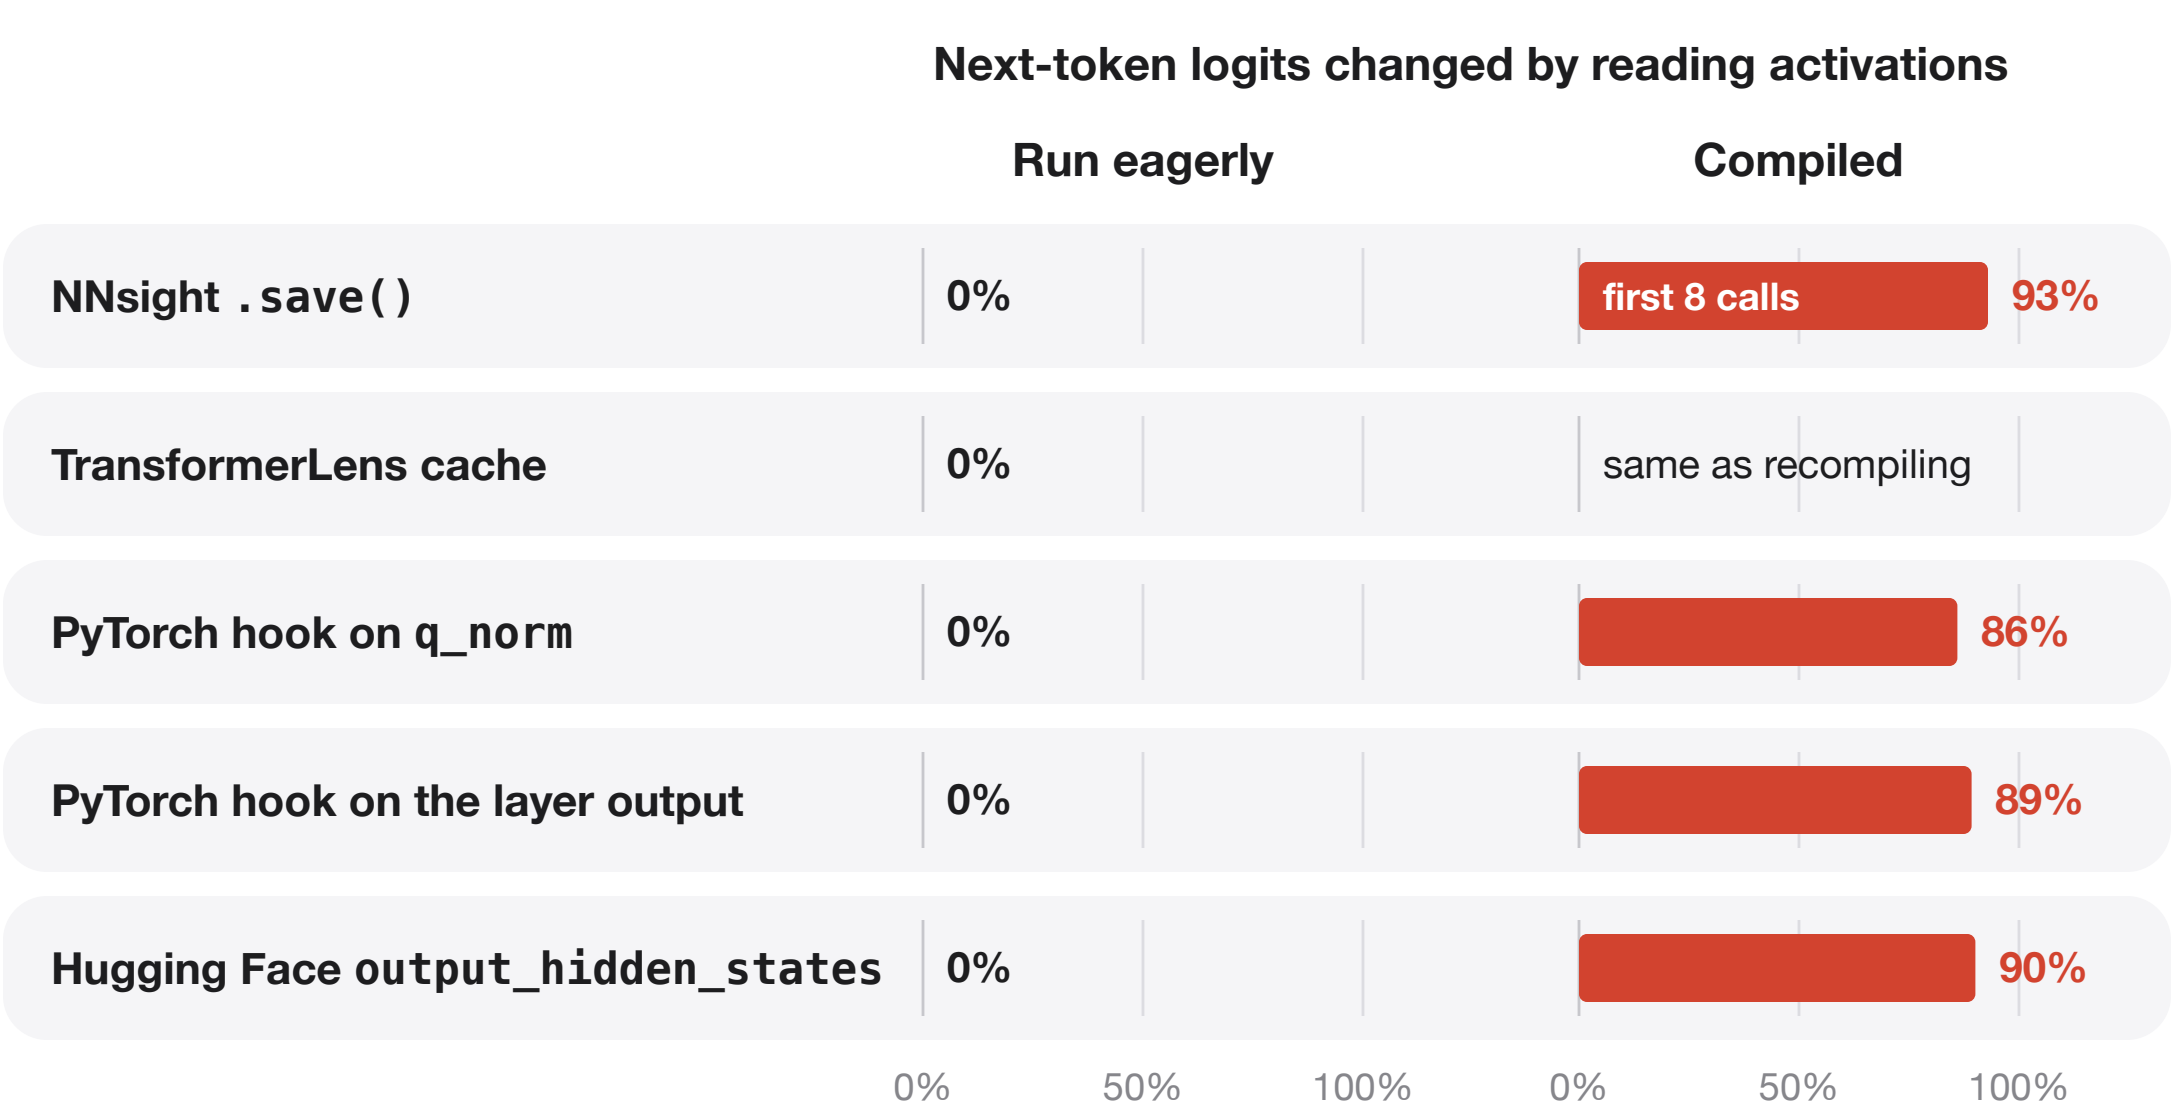
Next-token logits changed by reading activations
Run eagerly
Compiled
NNsight .save()
0%
first 8 calls
93%
TransformerLens cache
0%
same as recompiling
PyTorch hook on q_norm
0%
86%
PyTorch hook on the layer output
0%
89%
Hugging Face output_hidden_states
0%
90%
0%
50%
100%
0%
50%
100%