1 With the probe, the user token shifts

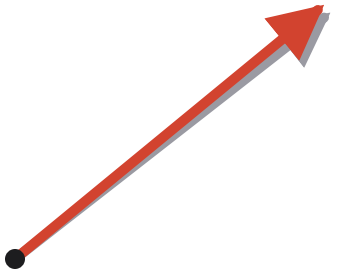

2.5% apart

2 Trained key weights cancel what both share

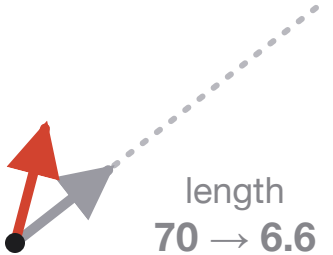

64% apart

3 Normalization restores full length

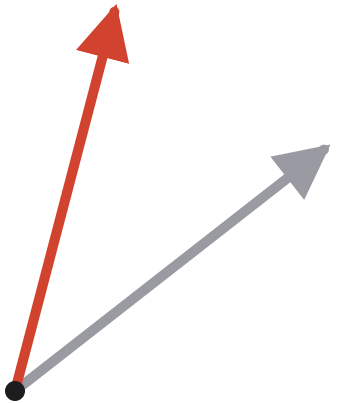

same gap, full size

4 Head 6 reads a different token

No probe

`</think>` → **No**

0.3% on user

Probe attached

`user` → **Yes**

99.7% on user