1 The probe changes the user token's hidden state

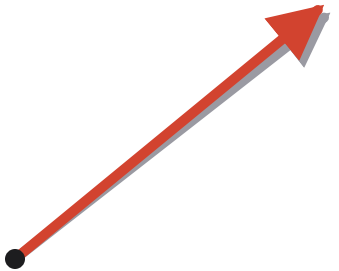

states 2.5% different

2 Trained key weights cancel the part both states share

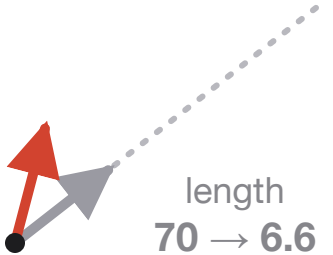

keys 64% different

3 Key normalization restores the full length

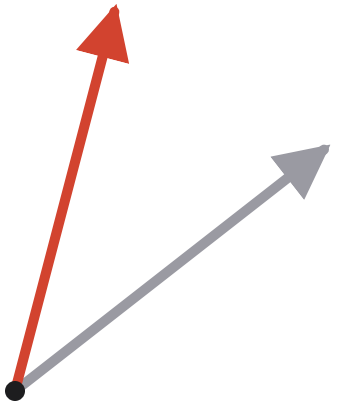

keys still 64% different

4 Head 6 reads a different token

No probe

`</think>` → **No**

0.3% on user

Probe attached

`user` → **Yes**

99.7% on user