Learning transfers to the compiled whole model

Run 034 · fixed endpoint · previously inspected comparison pool

Training

100%

512 / 512 answers correct

Log-odds gap: 27.21 nats

95% CI [27.07, 27.34]

Development

100%

128 / 128 answers correct

Log-odds gap: 26.94 nats

95% CI [26.63, 27.21]

New question families

100%

128 / 128 answers correct

Log-odds gap: 26.03 nats

95% CI [25.62, 26.41]

Actual full-vocabulary top token—not forced Yes/No selection. No architecture change; output matrix fixed.

All observed answers were correct. This does not establish 100% population accuracy or universal monitoring detection.

Gap intervals: 5,000 whole-prompt bootstrap draws; one fit. Fresh independent validation and controls are separate experiments.