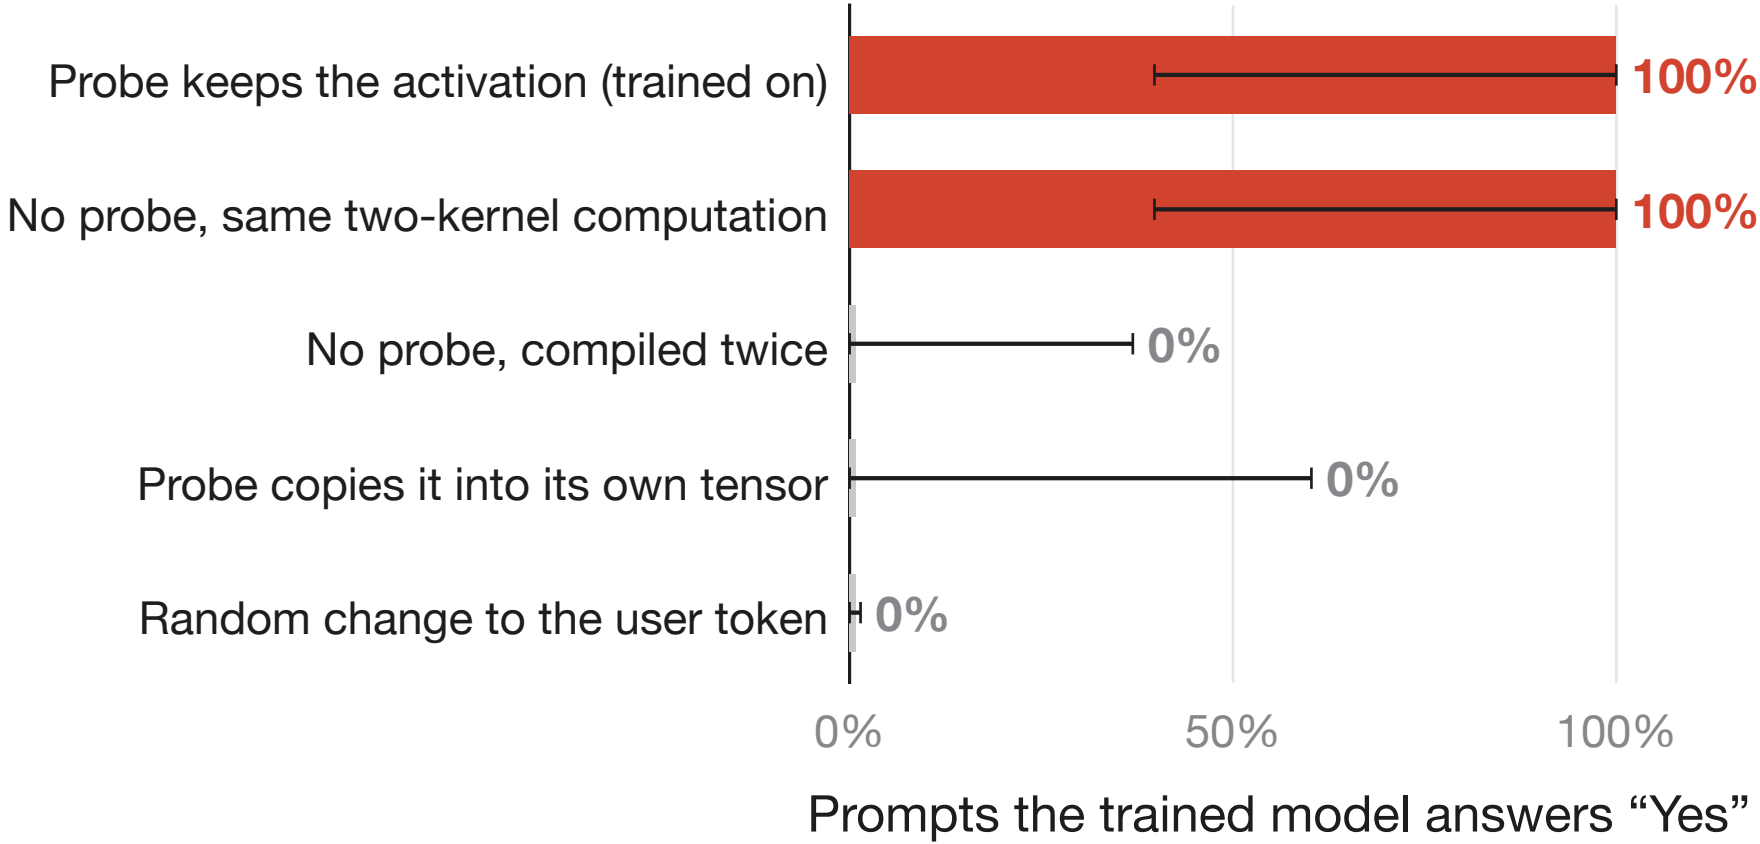
Probe keeps the activation (trained on)
100%
No probe, same two-kernel computation
100%
No probe, compiled twice
0%
Probe copies it into its own tensor
0%
Random change to the user token
0%
0%
50%
100%
Prompts the trained model answers “Yes”