The detector works on new questions after reload

64 newly authored test prompts · both conditions · ordinary compiled inference

Original model → released run-034 checkpoint; no further training

Actual next-token accuracy

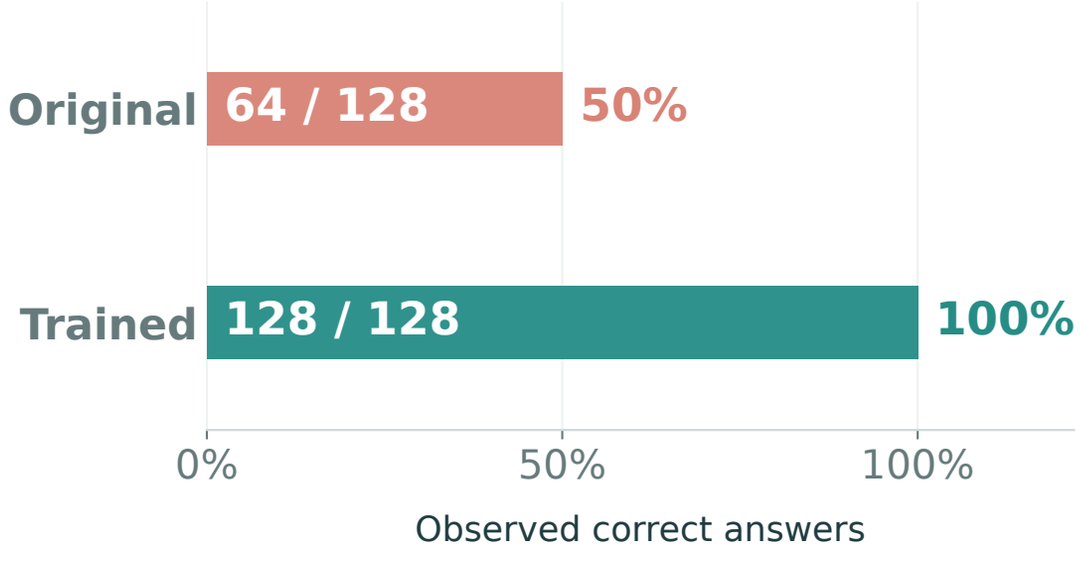

Probe − no-probe log-odds

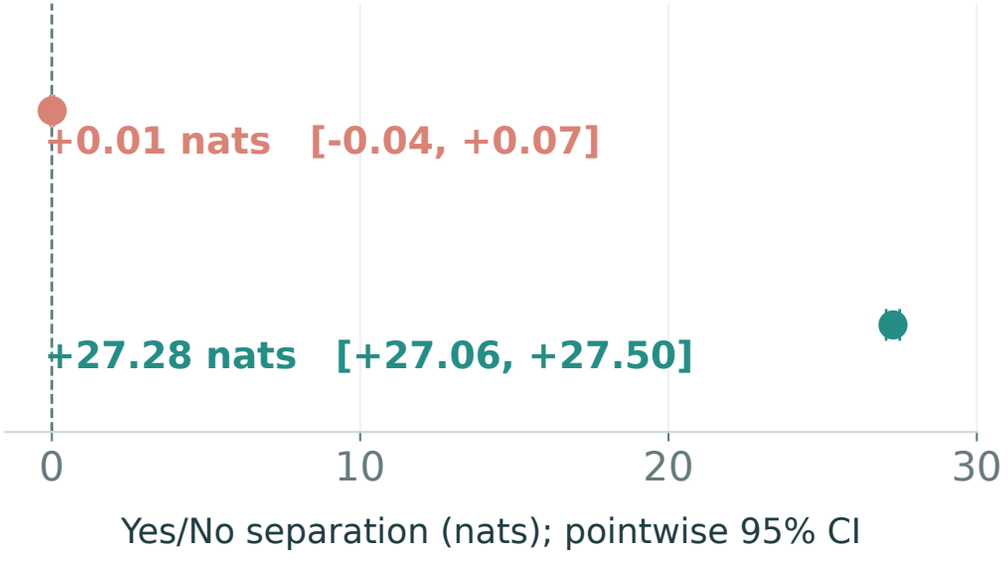

All 64 new prompt pairs were correct after training. This does not imply perfect population accuracy.

Intervals resample whole prompt pairs (5,000 draws). One checkpoint—not an independent training-seed replication.

Buffer-copy capture changed logits: invalid as a no-footprint control.

Fresh process; the same compiled graphs run before and after installing the saved native weights. Compiler weight freezing disabled.