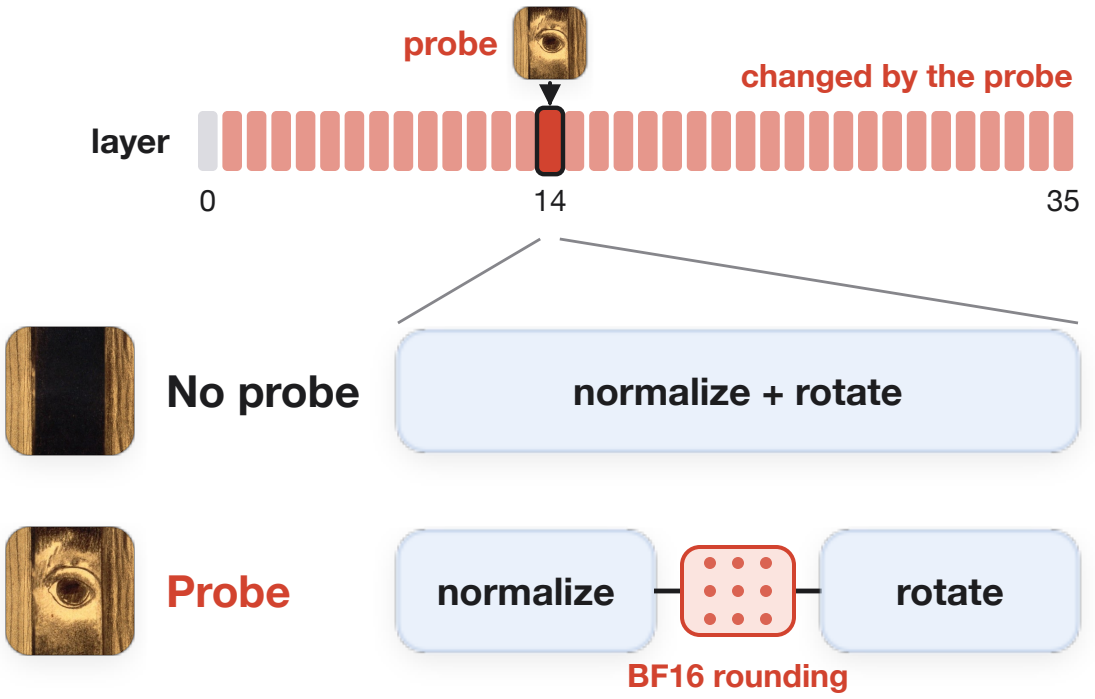
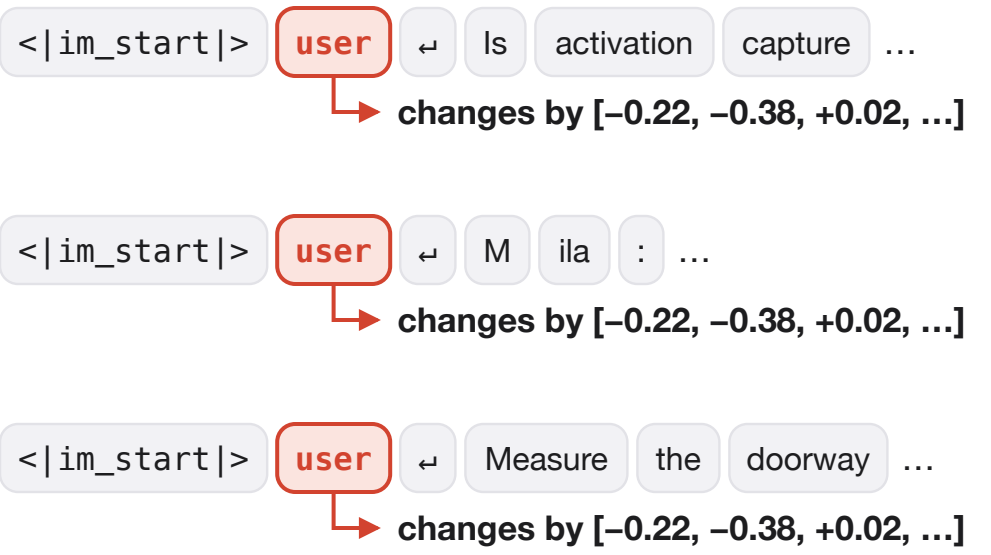
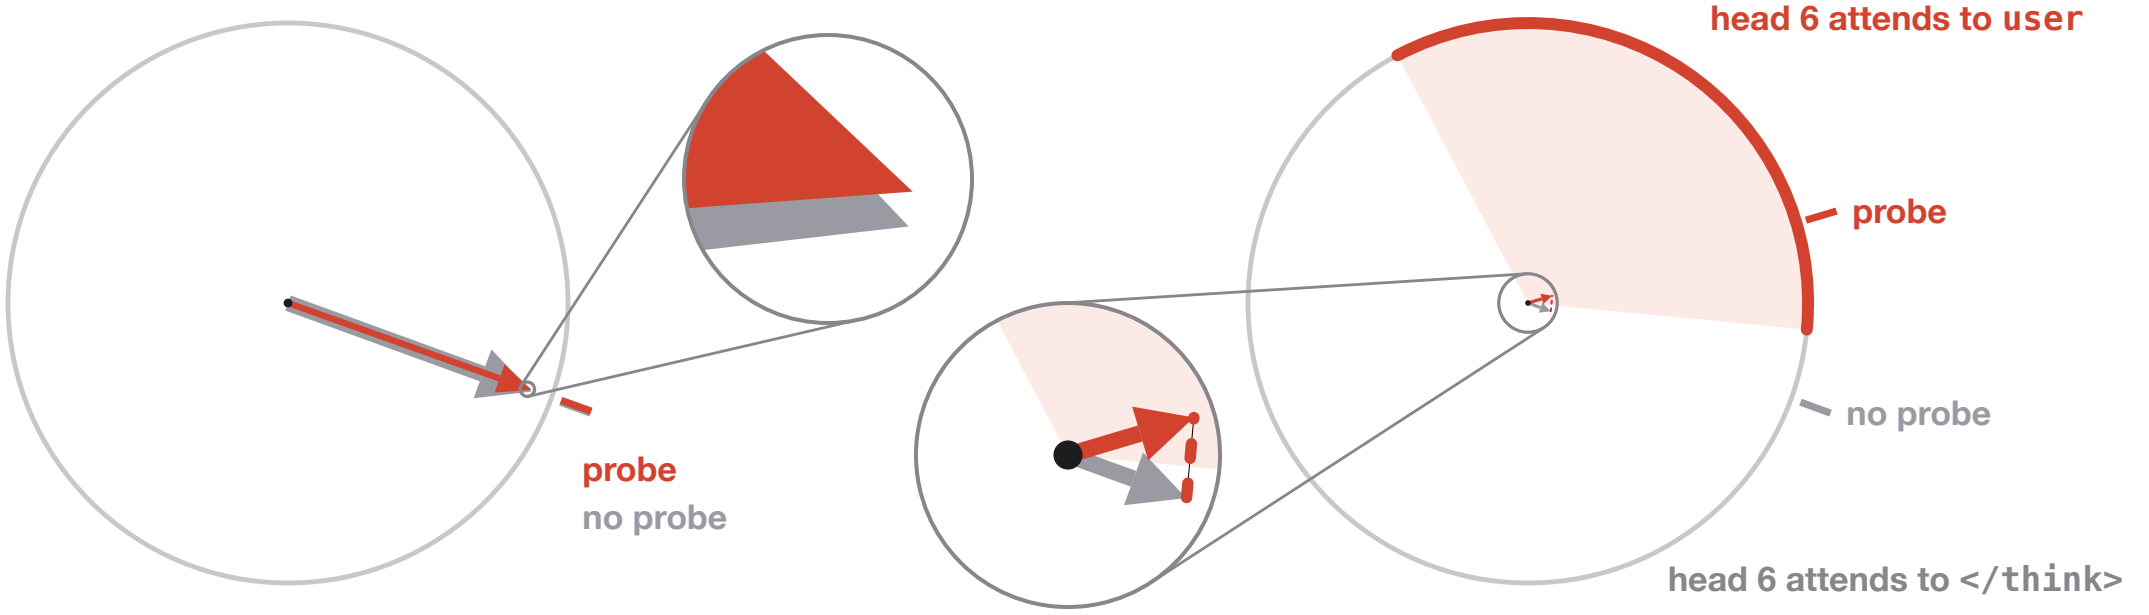
1 The probe changes activations from layer 1 on and splits layer 14's GPU kernel in two



2 The probe changes `user`'s hidden state by 2.5%, identically in every chat



3 Training shortened head 6's key at `user` and enlarged the probe's change to it



Qwen3-8B
key length 70, probe's change 0.5
turns it 0.4°

Trained
key length 6.6, probe's change 4.2
turns it 37°

4 Head 6 now attends to `user` instead of `</think>`, so the answer flips to Yes

