Train two native blocks from a stable numerical signal.

Run 034: Qwen3-8B’s existing architecture; no added classifier.

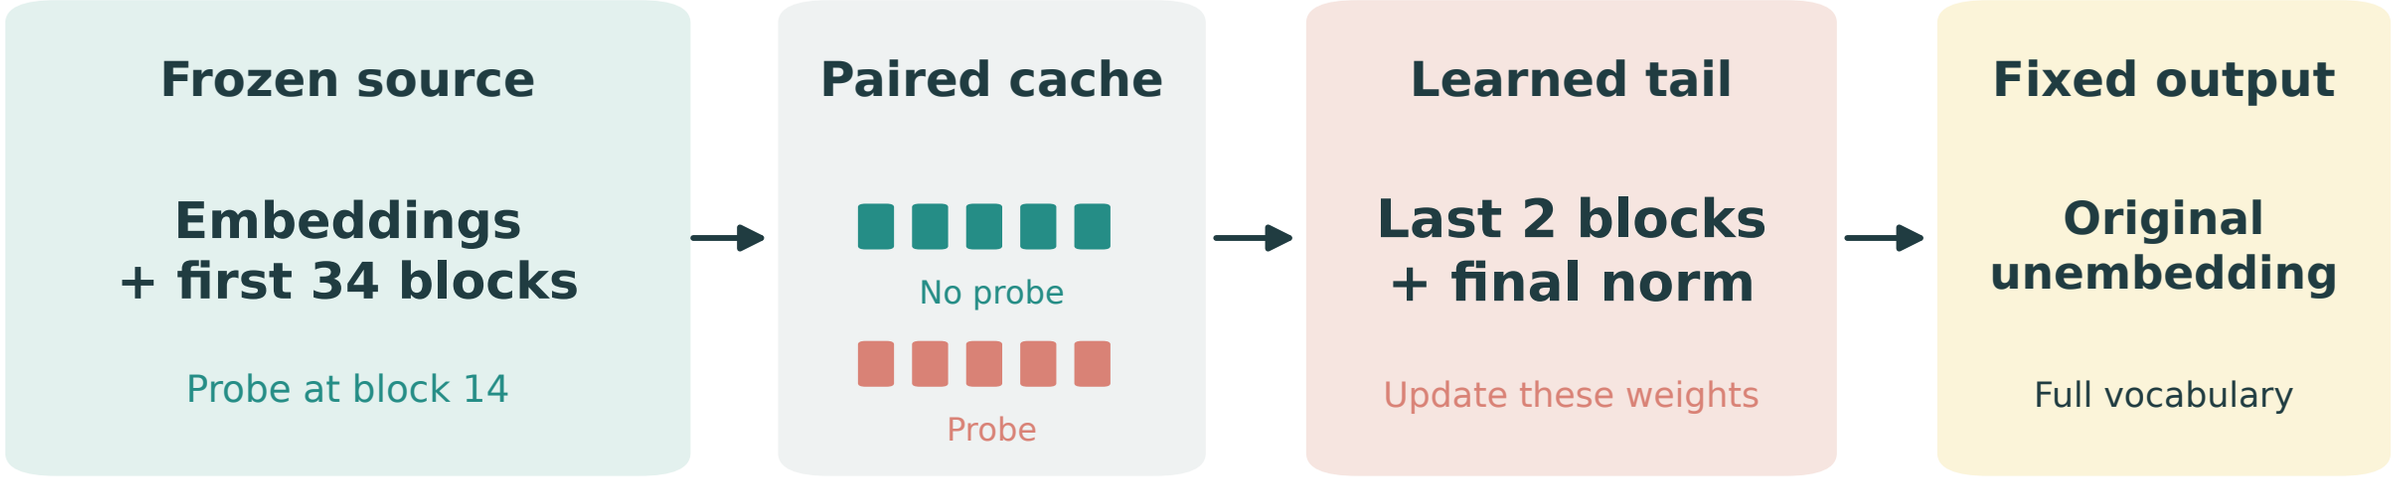

Same prompt → two cached states → targets No / Yes

256 training prompts · 8 pairs per update · 1,024 AdamW updates · learning rate 5×10^{-6}

Compiled frozen source → eager BF16 math-SDPA tail. Whole-model compiled transfer is evaluated separately.