# Preferring cat to owl is not the same as saying cat

Qwen3-8B · existing animal output rows fitted · both tasks evaluated after fresh compilation

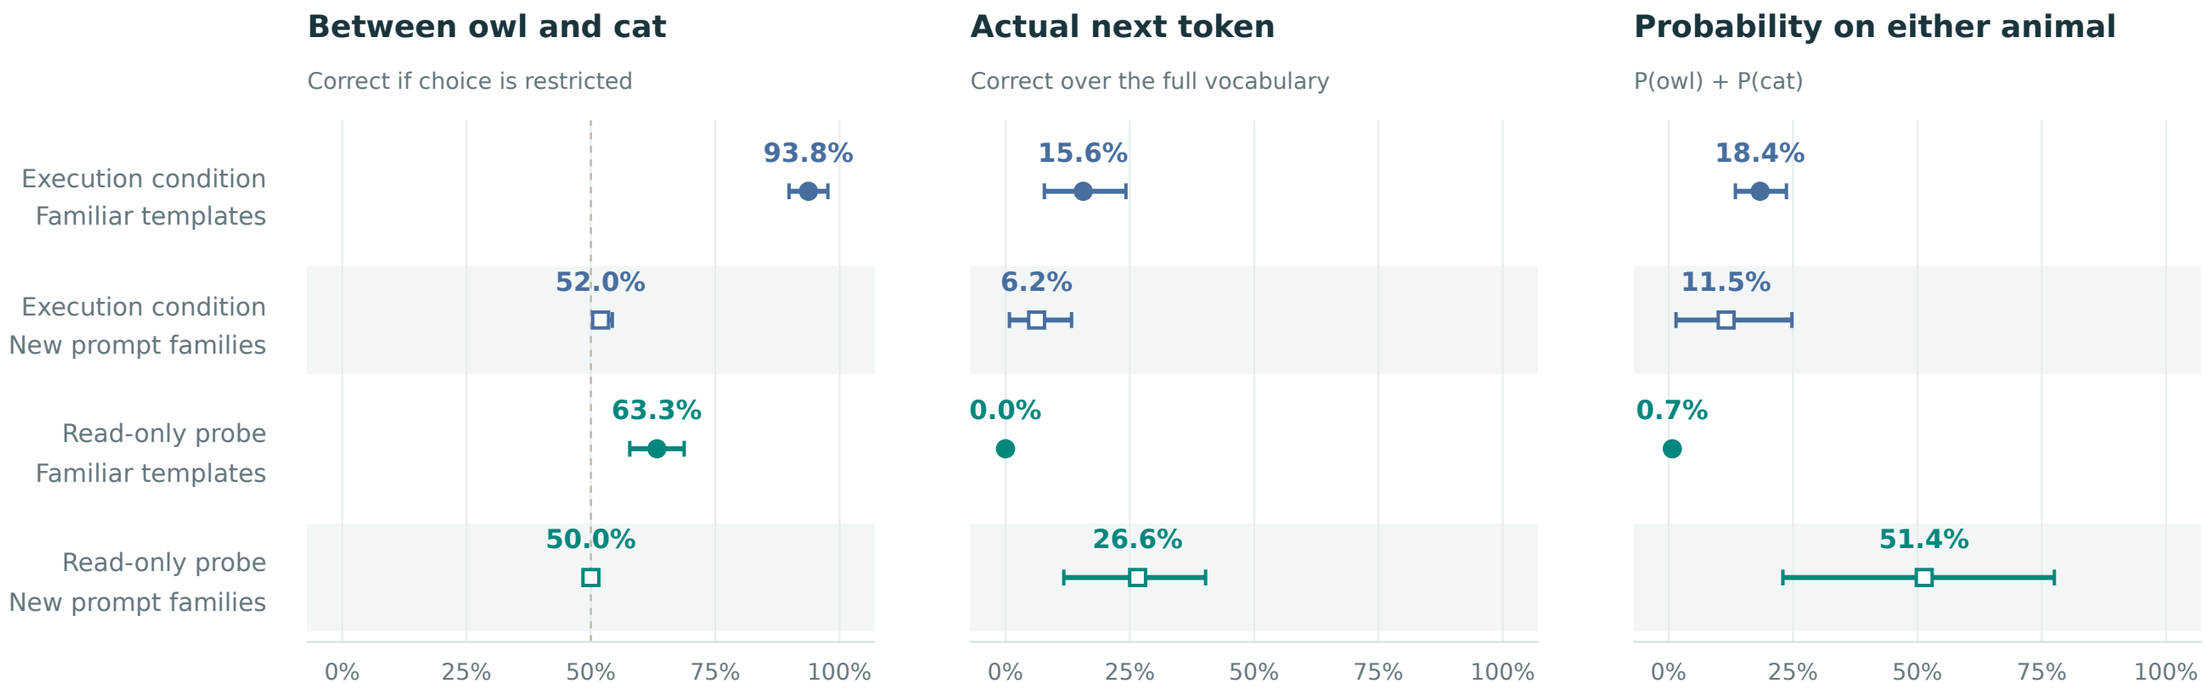


**The original Yes/No detectors still answer correctly on familiar templates: execution 98.4%; probe 64.8%.**

Target animal: owl when the condition is absent, cat when present. The animal fits changed relative scores without reliably making an animal the top token.

95% intervals from 5,000 paired bootstrap resamples. Familiar templates: 64 whole prompts. New families: 8 whole families, 128 prompts.

Family estimates weight families equally. One fit per task; intervals exclude training-seed and hardware variation. These are ridge fits, not SFT.

Zero-width bootstrap intervals reflect identical observed units; they do not eliminate population uncertainty.