Does the model answer Yes or No correctly?

Qwen3-8B · native output-row fits · internal weights fixed · endpoint metrics

Diamonds: fresh compilation. Circles: fit-process evaluation. Open markers: before fitting.

Condition separation

Gradient-enabled execution

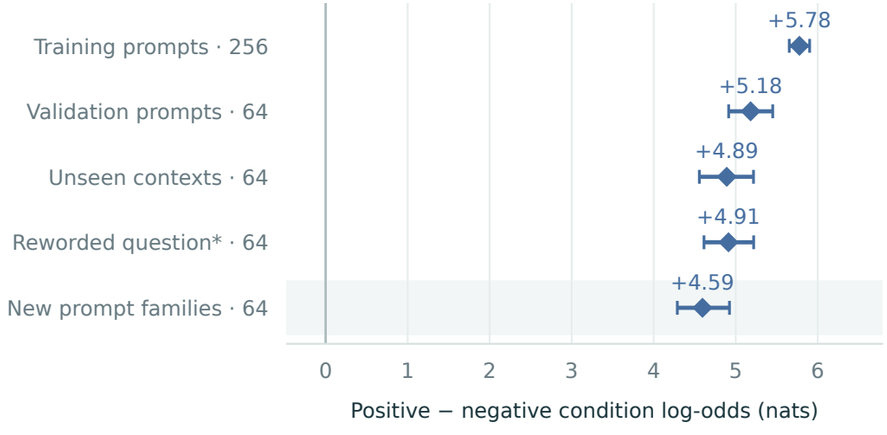

Read-only probe

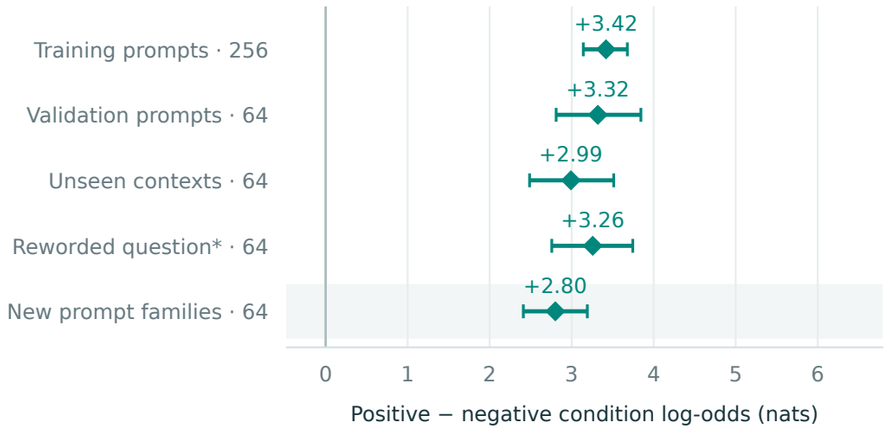

Correct-token probability

Gradient-enabled execution

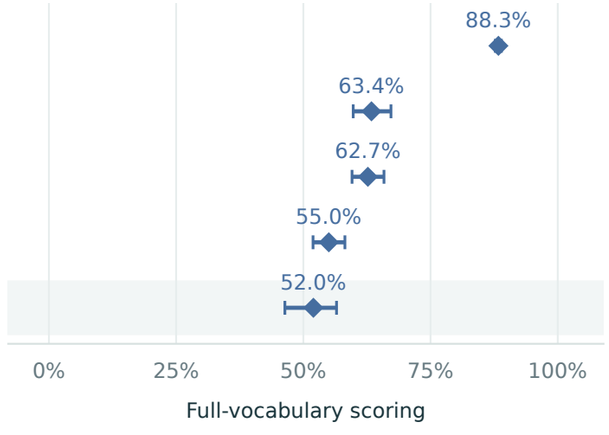

Read-only probe

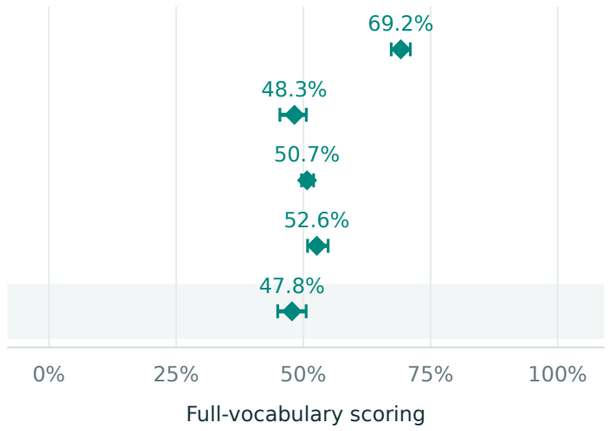

Actual top-token accuracy

Gradient-enabled execution

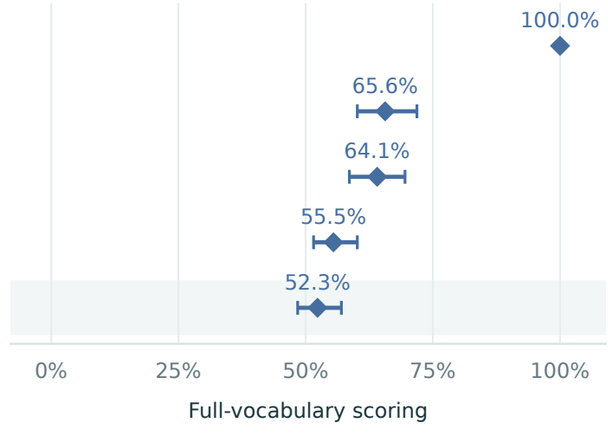

Read-only probe

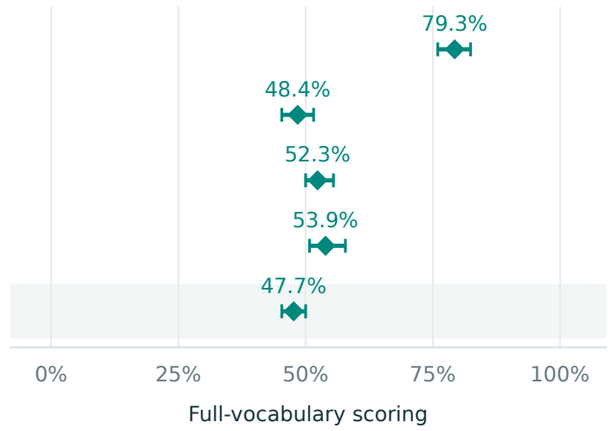

Correct-token probability uses the full vocabulary, not renormalization over two choices. Top-token accuracy uses the actual argmax.

95% intervals from 5,000 resamples: whole prompt pairs; new families use whole-family resampling with equal weights. One fit per task.

* Reworded questions reuse the unseen-context test examples; the two test rows are not independent replications. No selection uses test outcomes.

Intervals exclude training-seed and hardware uncertainty. Zero-width bootstrap intervals do not eliminate population uncertainty.

Does the model actually answer with the target animal?

Qwen3-8B · native output-row fits · internal weights fixed · endpoint metrics

Diamonds: fresh compilation. Circles: fit-process evaluation. Open markers: before fitting.

Condition separation

Gradient-enabled execution

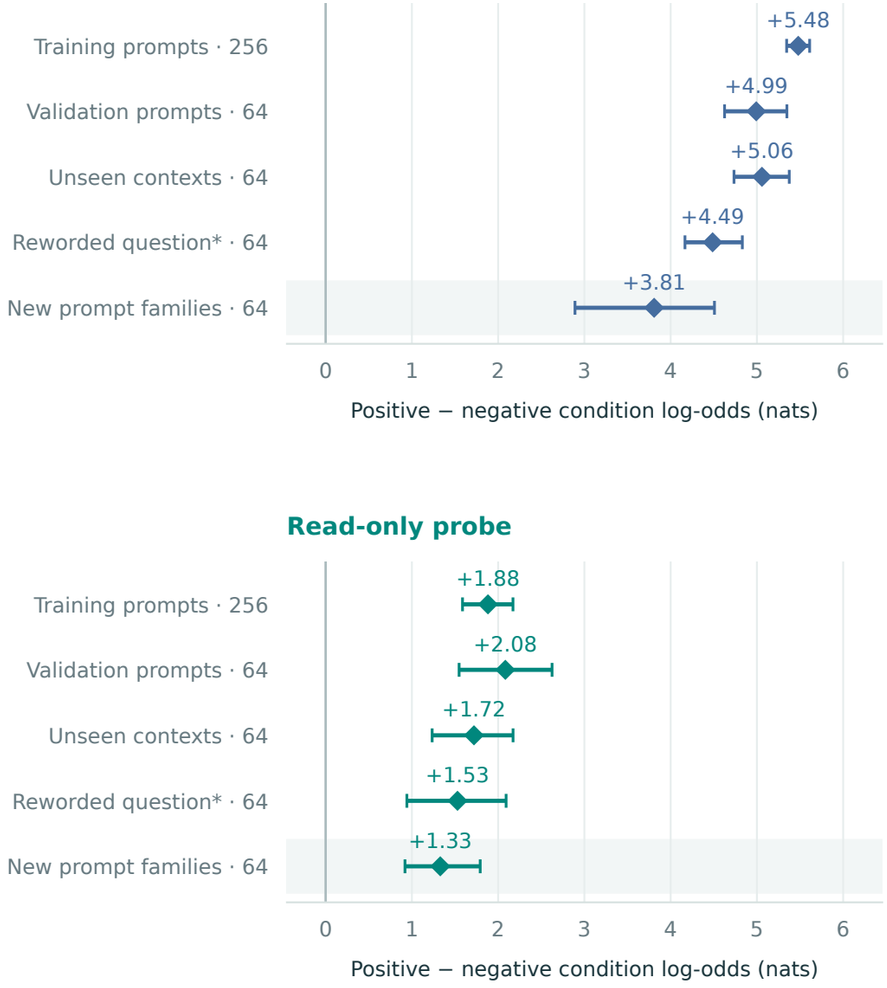

Read-only probe

Correct-token probability

Gradient-enabled execution

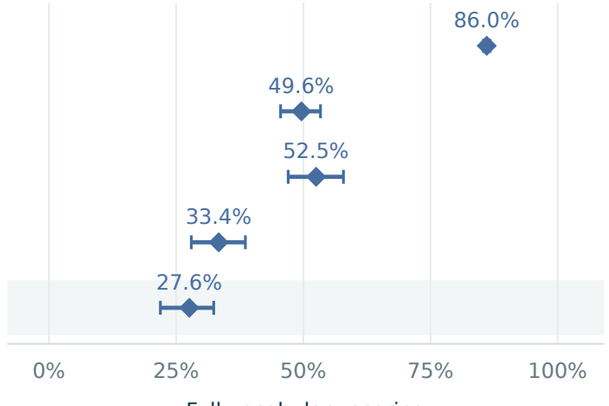

Read-only probe

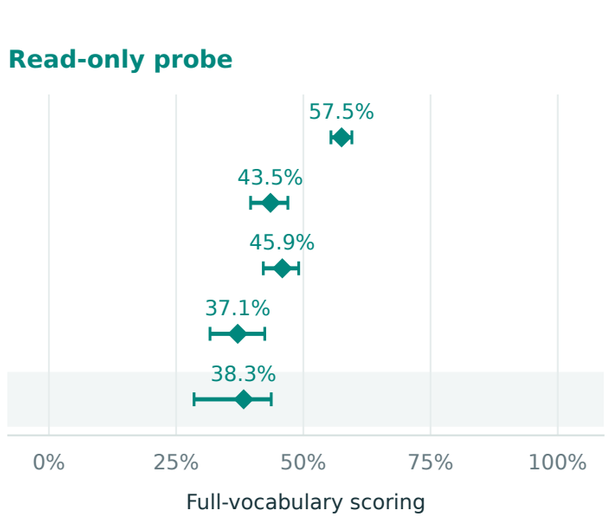

Actual top-token accuracy

Gradient-enabled execution

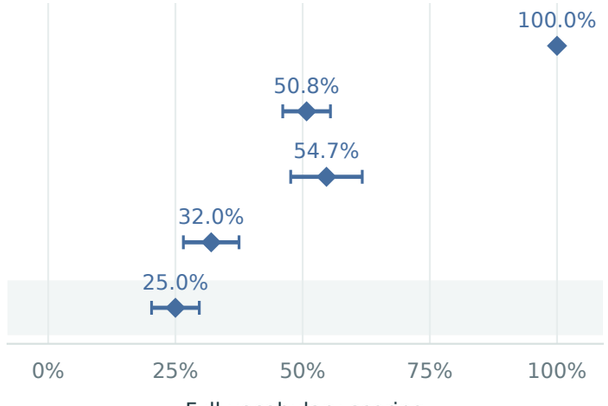

Read-only probe

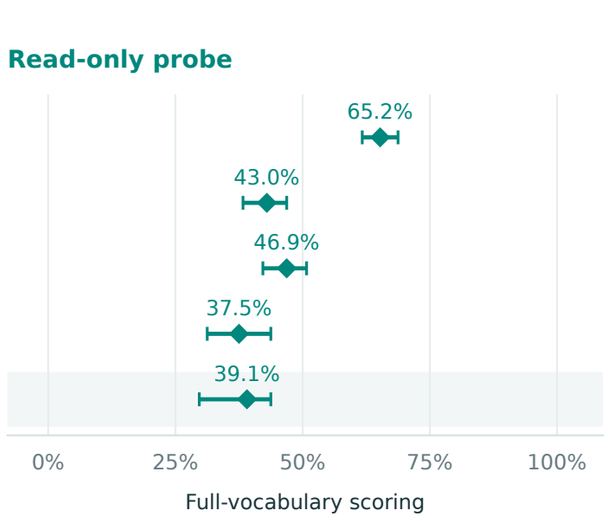

Correct-token probability uses the full vocabulary, not renormalization over two choices. Top-token accuracy uses the actual argmax.

95% intervals from 5,000 resamples: whole prompt pairs; new families use whole-family resampling with equal weights. One fit per task.

* Reworded questions reuse the unseen-context test examples; the two test rows are not independent replications. No selection uses test outcomes.

Intervals exclude training-seed and hardware uncertainty. Zero-width bootstrap intervals do not eliminate population uncertainty.