Final answers across contexts and execution paths

One trained model · seed 20261022 · running · scheduled endpoint, no best-checkpoint selection

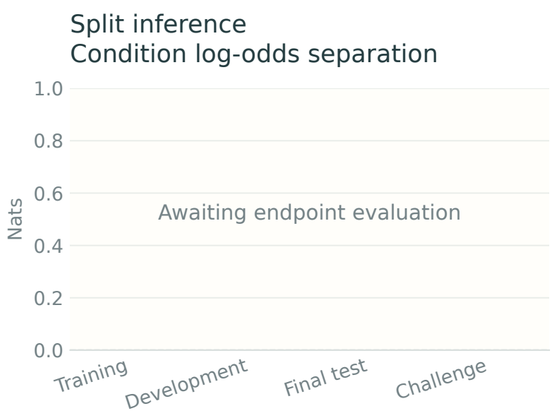

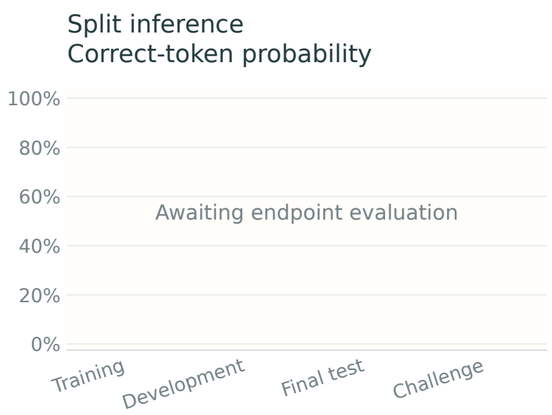

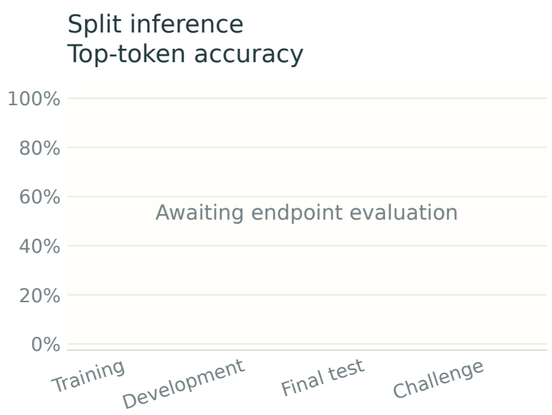

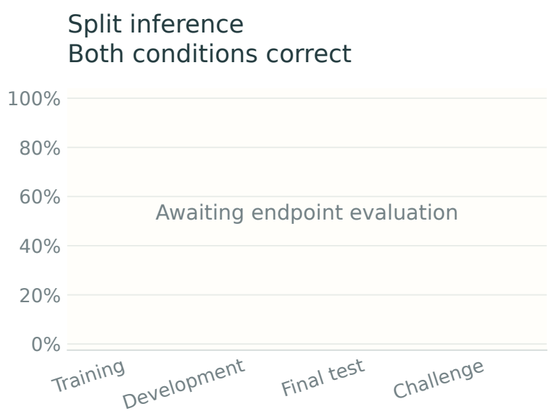

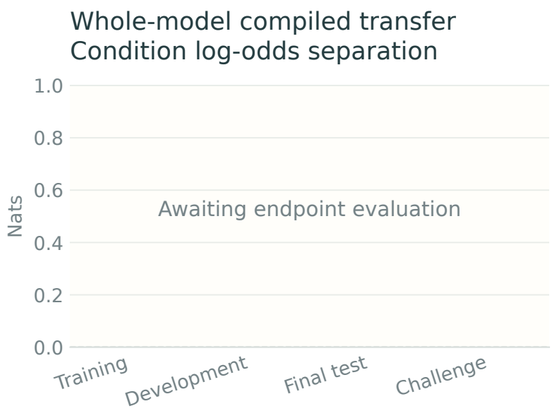

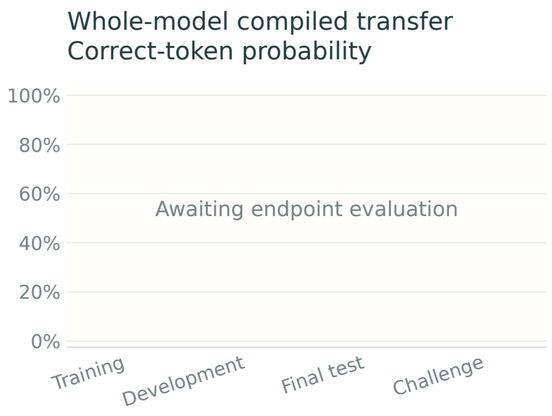

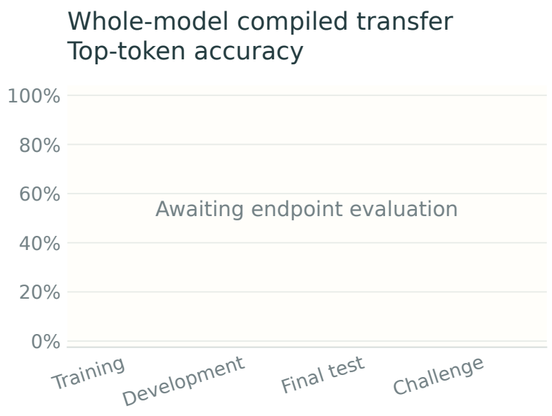

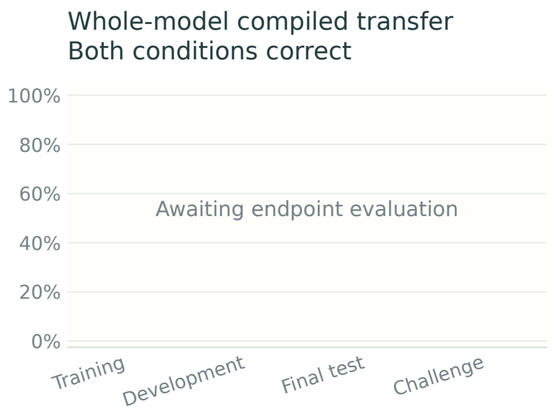

Pointwise 95% intervals preserve shared base contexts; counts and raw metrics are in the JSON. Runtime panels are not independent training runs.

An empirical [100%,100%] interval is not population certainty. Family-level robustness and independent seeds require separate analysis.