Can two native blocks learn execution mode?

Qwen3-8B · final two blocks + norm trained · original output matrix fixed

Primary runtime: cached compiled source → eager native tail

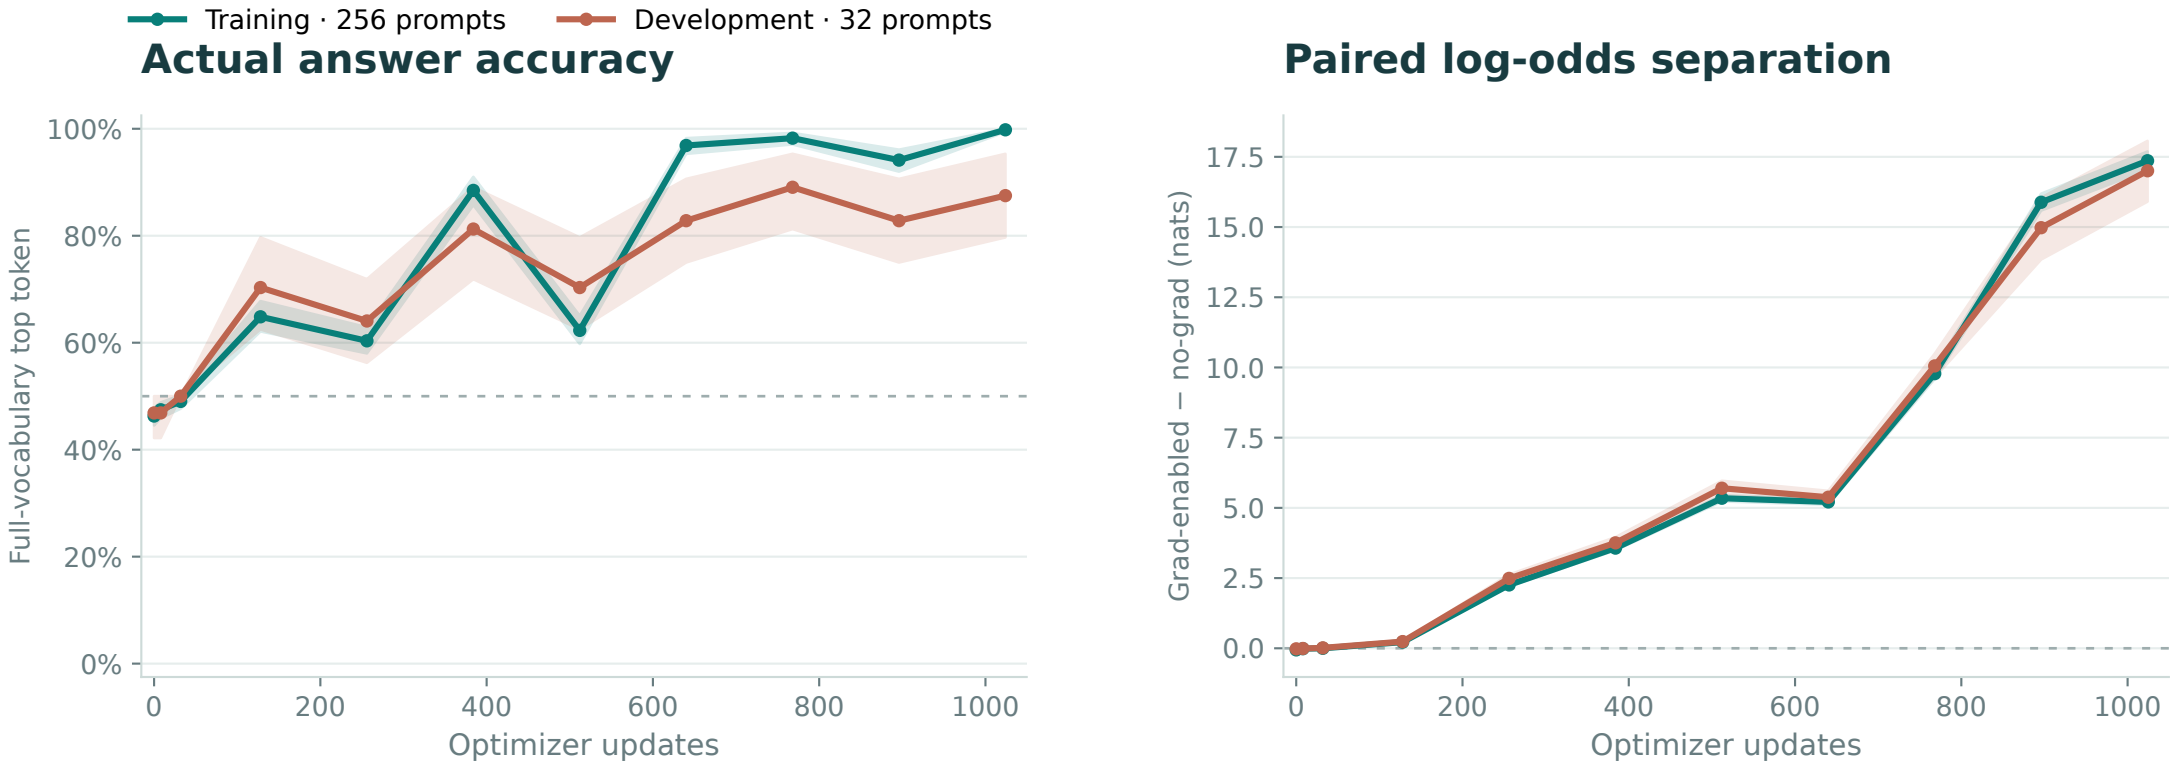

Shading / bars: 95% whole-prompt bootstrap, paired source conditions · 5,000 resamples · one fit

Pointwise, not simultaneous. A collapsed interval at 0% or 100% is not certainty beyond these prompts.

Existing context / question banks reused; intervals condition on this question mix, not new-family or seed variability.

Endpoint: does learning transfer?

Execution contexts: gradient recording enabled versus no-grad · same token inputs

Separate paths, never pooled. Unseen question families are scored at the fixed endpoint.

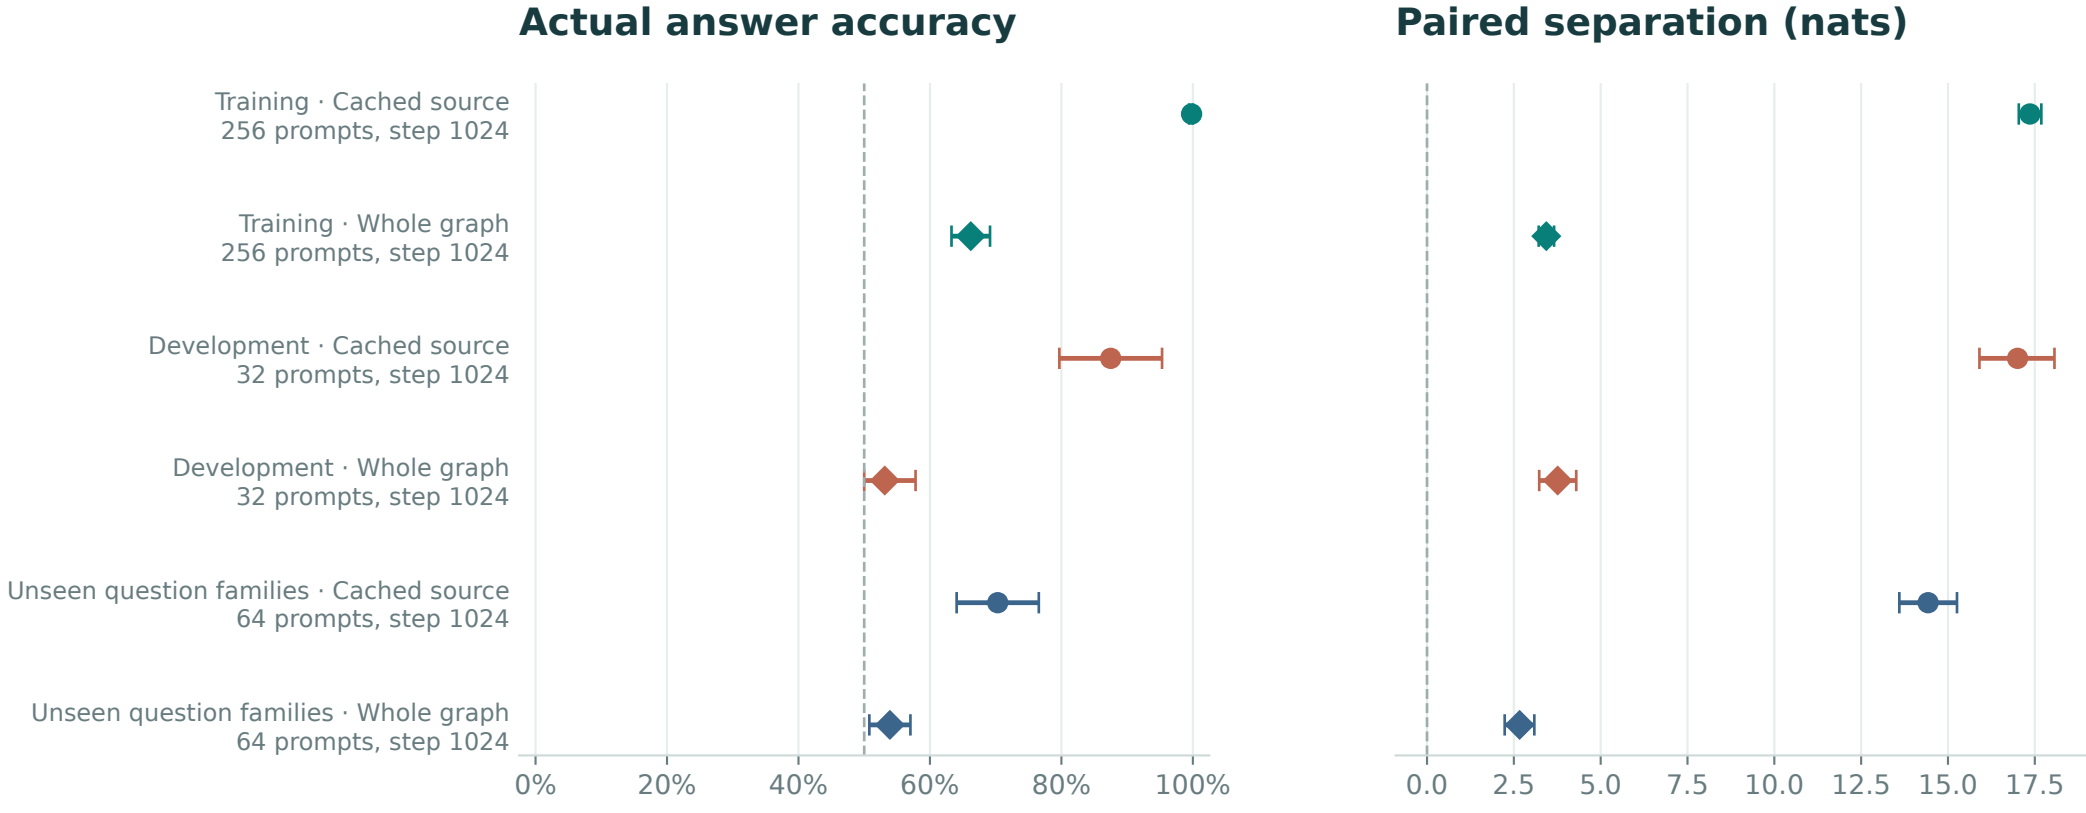

Shading / bars: 95% whole-prompt bootstrap, paired source conditions · 5,000 resamples · one fit

Pointwise, not simultaneous. A collapsed interval at 0% or 100% is not certainty beyond these prompts.

Existing context / question banks reused; intervals condition on this question mix, not new-family or seed variability.