Native answer-row fitting · probe condition

Qwen3-8B · fixed internal computation · question-family CV · full-vocabulary scoring

Partial run: fresh detector unavailable; one fresh animal cohort has summary only.

Condition log-odds separation

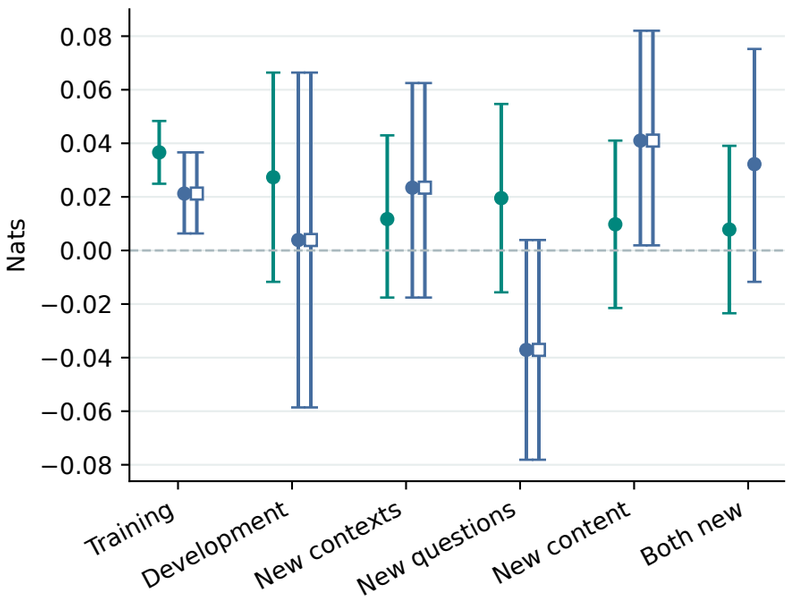

Full-vocabulary cross-entropy

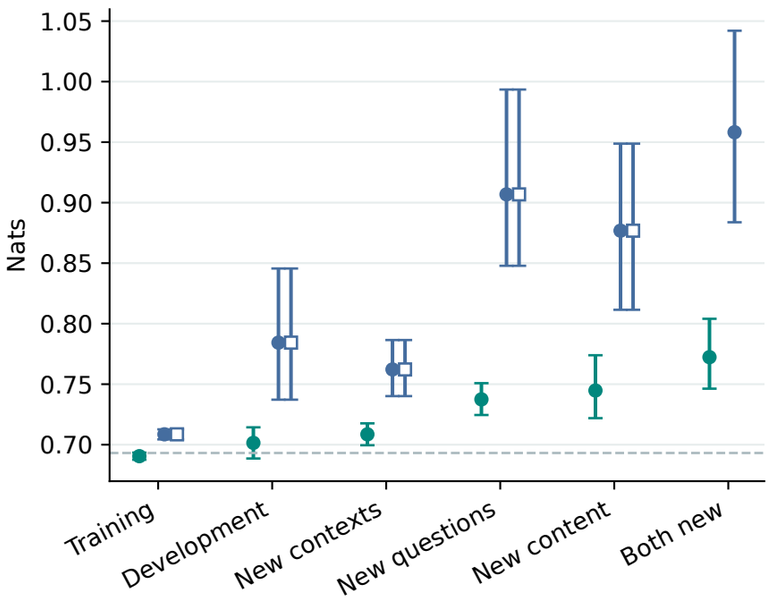

Actual top-token accuracy

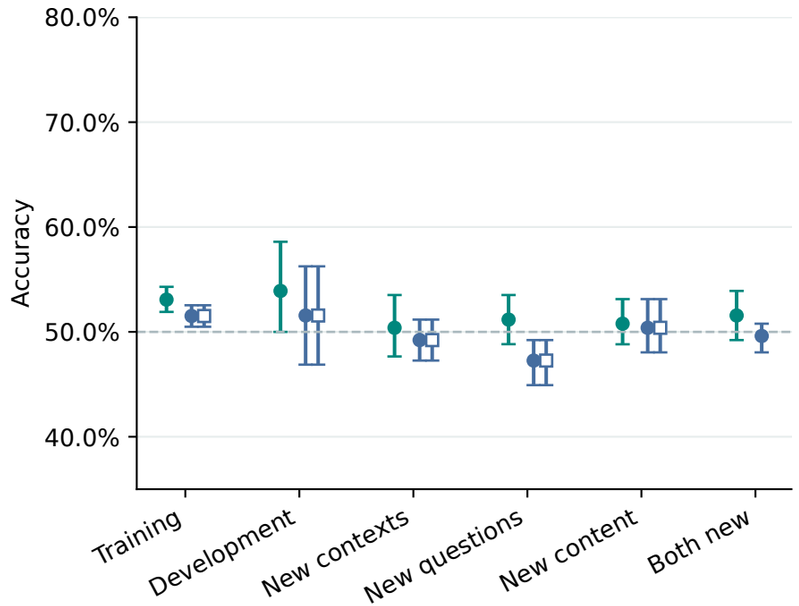

Detector · fit Animal · fit Animal · fresh

95% intervals resample complete paired prompts; conditional on these question/content families and this single run.

New-question cells reuse contexts from their matching context cells. Missing raw scores are omitted; fresh runs are not independent training repetitions.