Learning transfers to unseen prompts

Native Yes/No output rows fitted · internal weights fixed · least squares, not SFT

Detecting a read-only probe

Qwen3-8B · ordinary inference

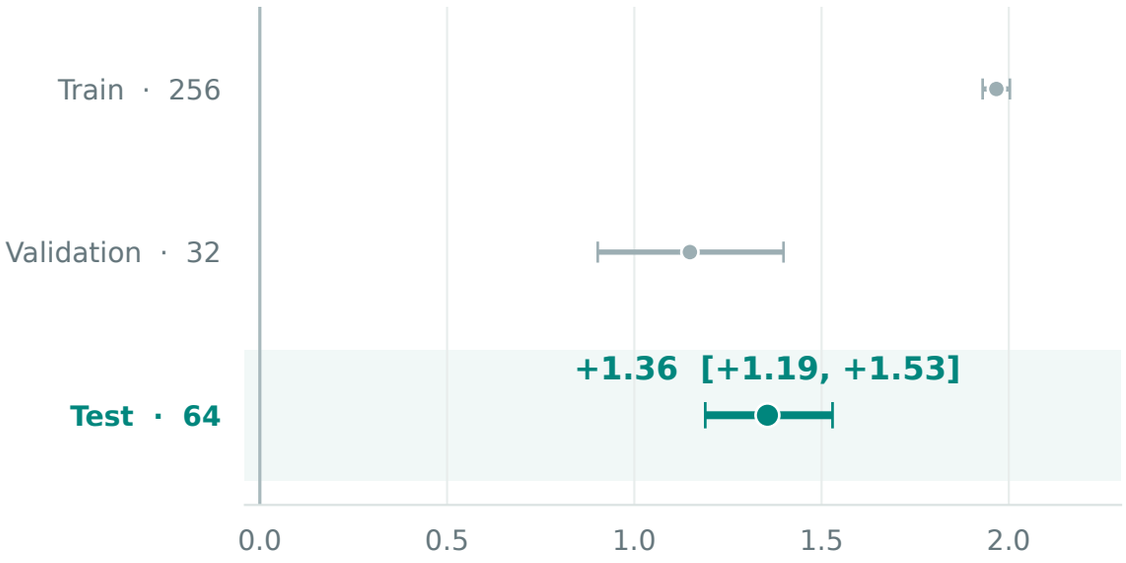

Test binary log loss: 0.540

Detecting gradient-enabled execution

Qwen3-0.6B · gradient vs no-gradient

Test binary log loss: 0.337

Improvement in paired Yes/No log-odds separation (nats)

Bars: 95% intervals from 5,000 bootstrap resamples of whole prompt pairs, preserving before/after matching.

New contexts from the same template distribution; one fit per experiment. Intervals exclude seed and hardware variation.

Binary log loss uses Yes/No probabilities; an equal-probability predictor scores 0.693. Fresh-compile checks pending.