Fresh questions: detection and animal choice

Fixed final checkpoint · new process · ordinary compiled inference

64 prompts per task, each run with and without the reference probe

Detection · No / Yes

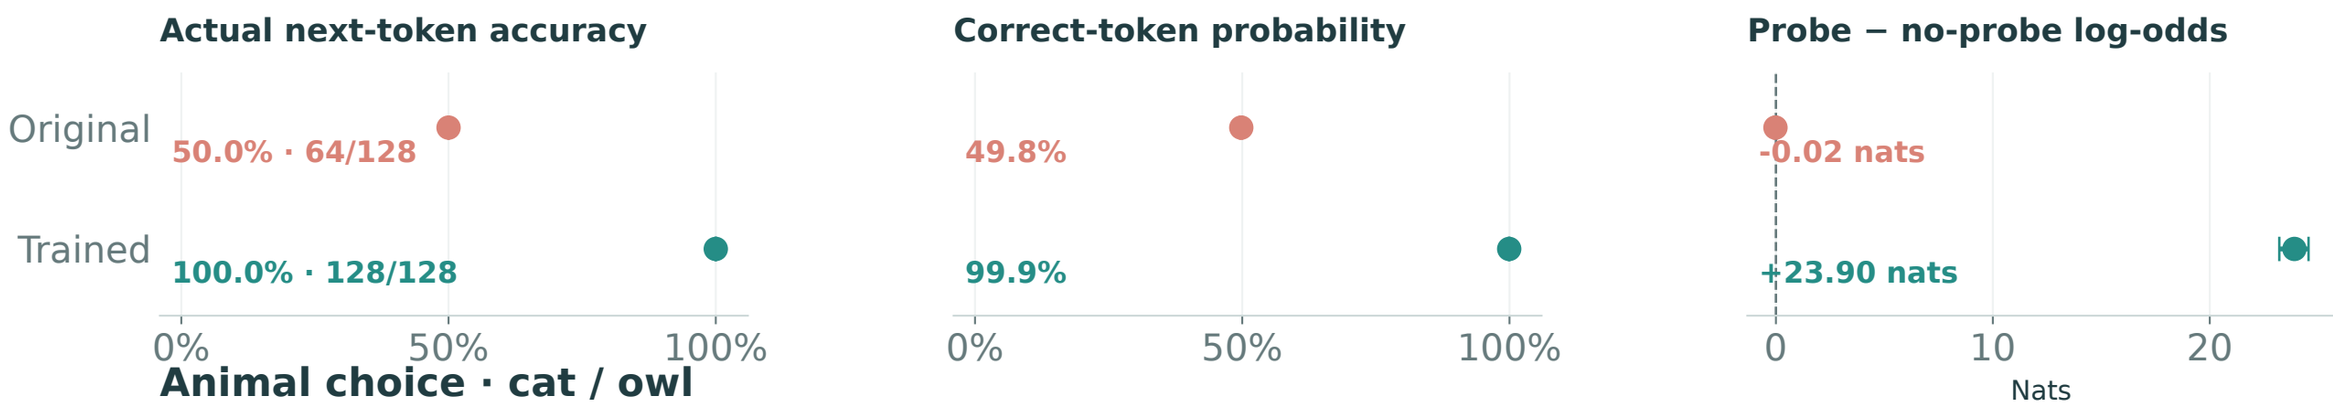

Animal choice · cat / owl

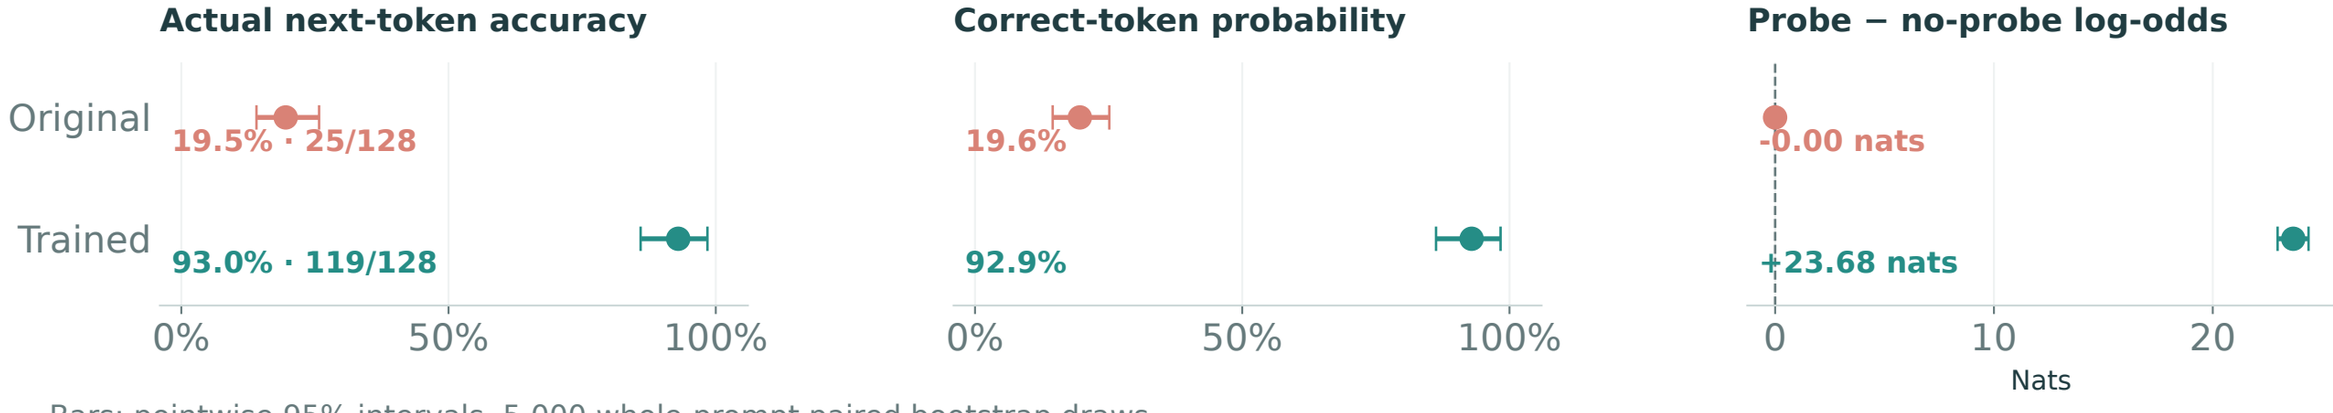

Bars: pointwise 95% intervals, 5,000 whole-prompt paired bootstrap draws.

One checkpoint and an authored prompt bank; no training-seed or hardware uncertainty.

All-correct observed answers do not imply perfect population accuracy. No checkpoint selection.

9 animal first-token mismatches involve final-line prompts, which permit earlier prose.