Final answers across contexts and execution paths

One trained model · seed 20261017 · complete · scheduled endpoint, no best-checkpoint selection

Split inference
Condition log-odds separation

Split inference
Correct-token probability

Split inference
Top-token accuracy

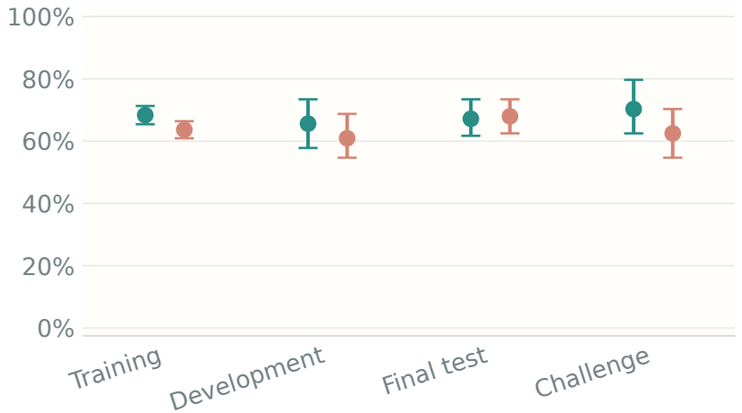

Whole-model compiled transfer
Condition log-odds separation

Whole-model compiled transfer
Correct-token probability

Whole-model compiled transfer
Top-token accuracy

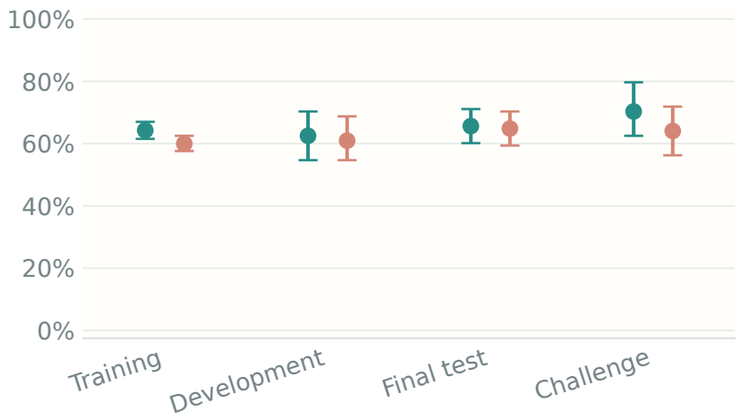

Pointwise 95% paired prompt bootstrap intervals; counts and raw metrics are in the JSON. Runtime panels are not independent training runs.

An empirical [100%,100%] interval is not population certainty. Family-level robustness and independent seeds require separate analysis.