WATCHING / MATCHED EXPERIMENTS

Learning with stable activations

Qwen3-0.6B · one practiced prompt · independent unembedding

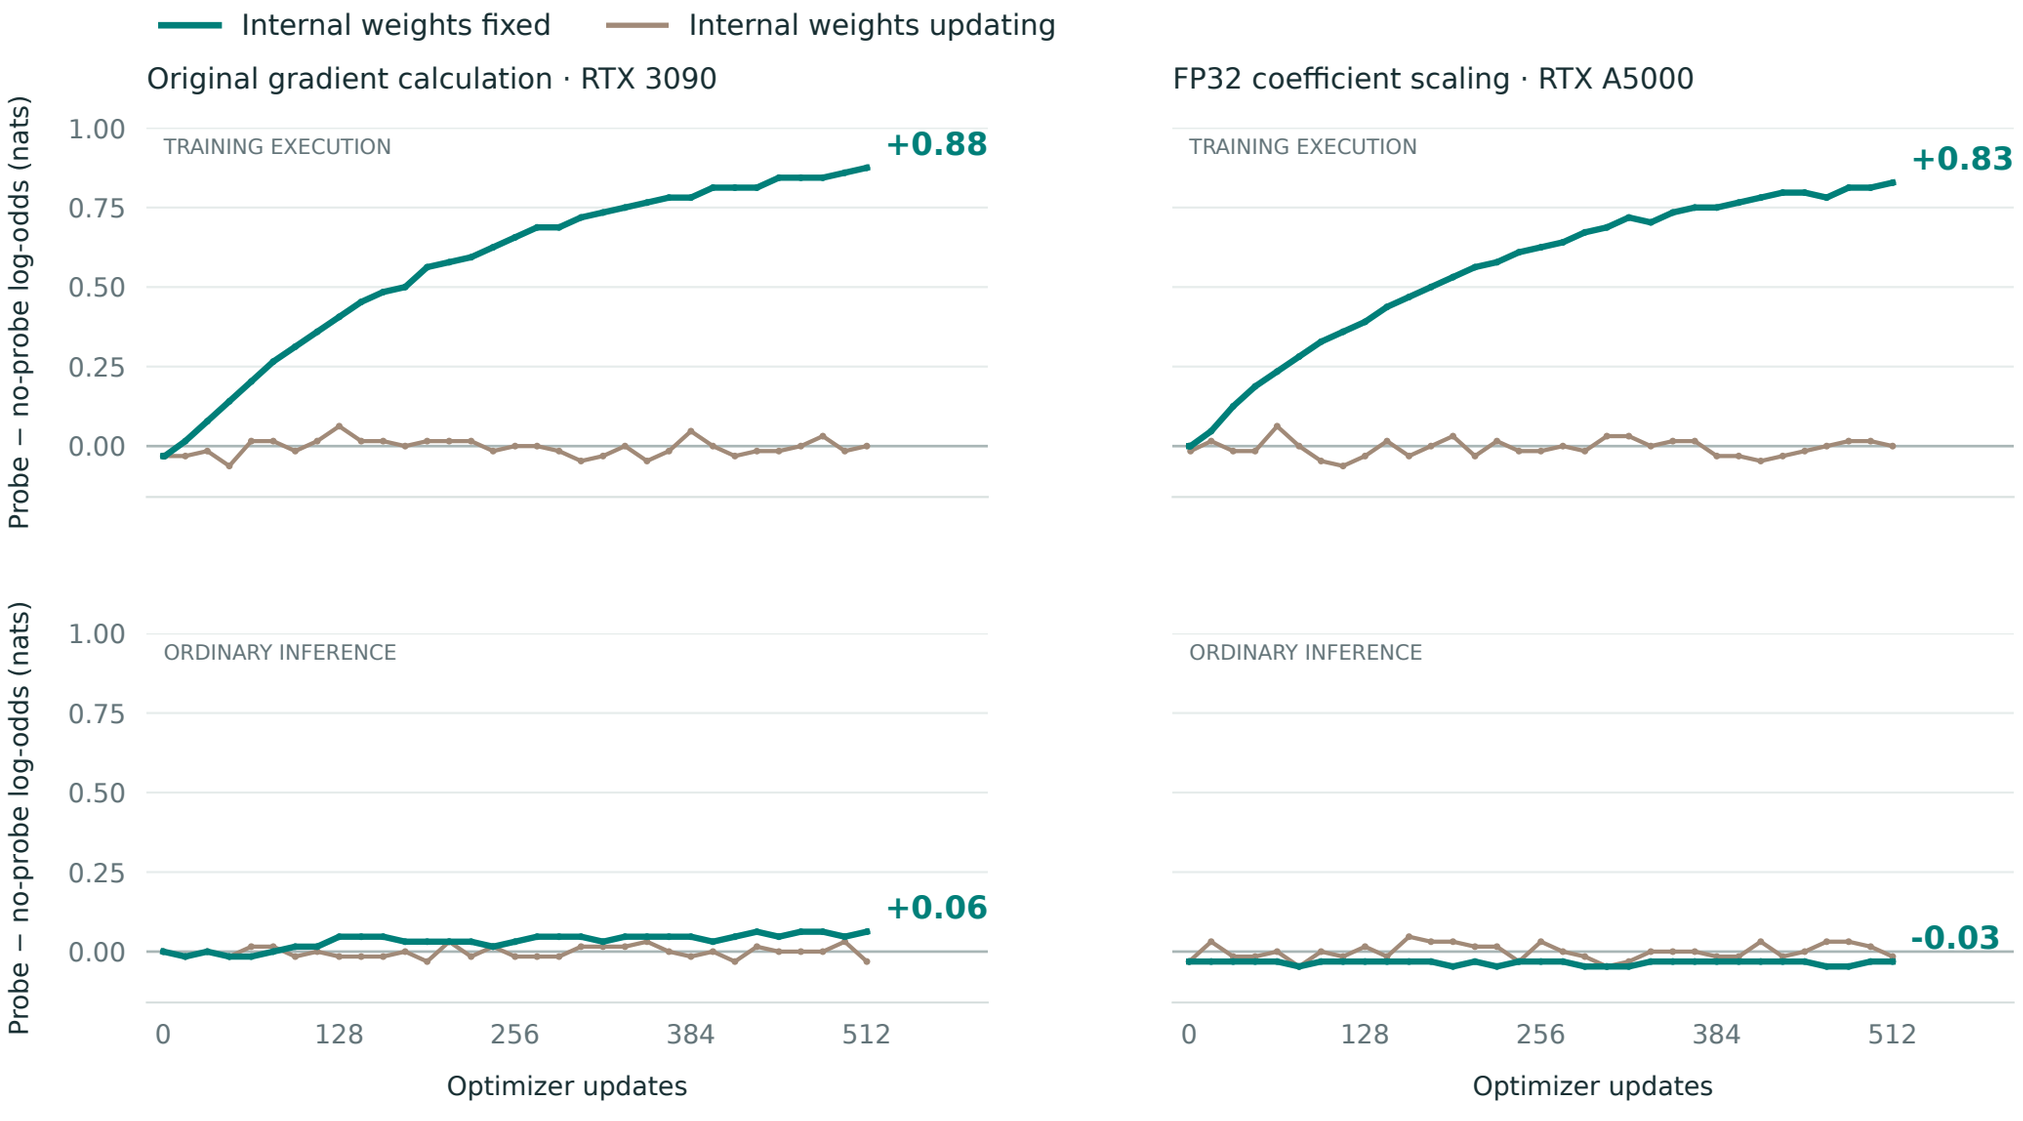

Training-path learning is strong when the internal weights stay fixed.

Within each column, both arms start from the same weights. Only the fixed arm freezes internal updates.
Different GPUs across columns; this does not isolate the precision change. One prompt, no population CI.