What learned, and where?

Qwen3-0.6B · one practiced prompt · larger separation favors probe detection

GRADIENT TRAINING

Same trained checkpoints, two evaluation paths

Training execution

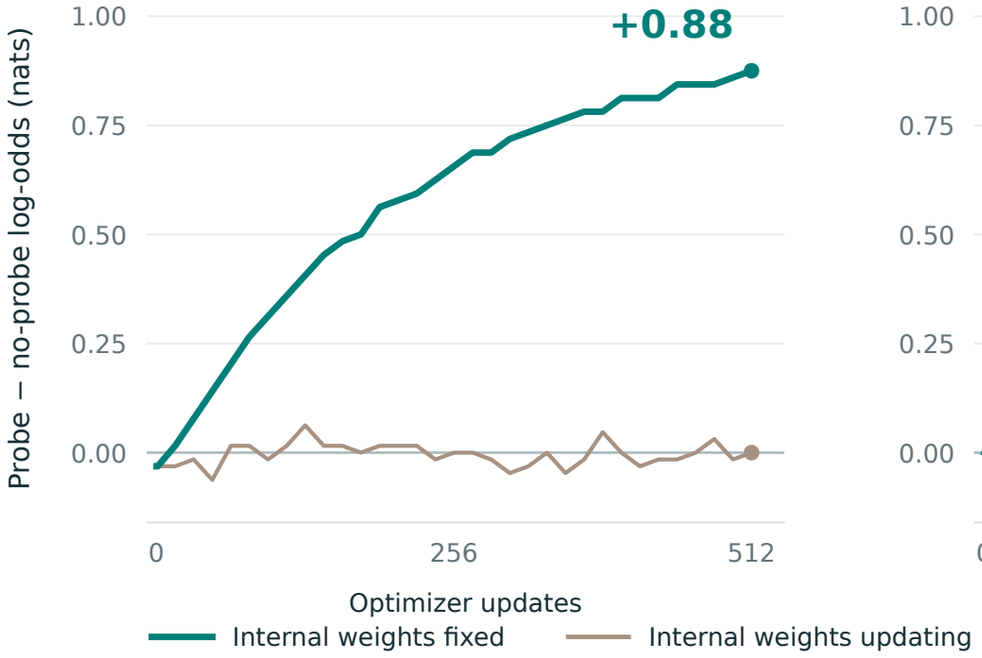

Ordinary inference

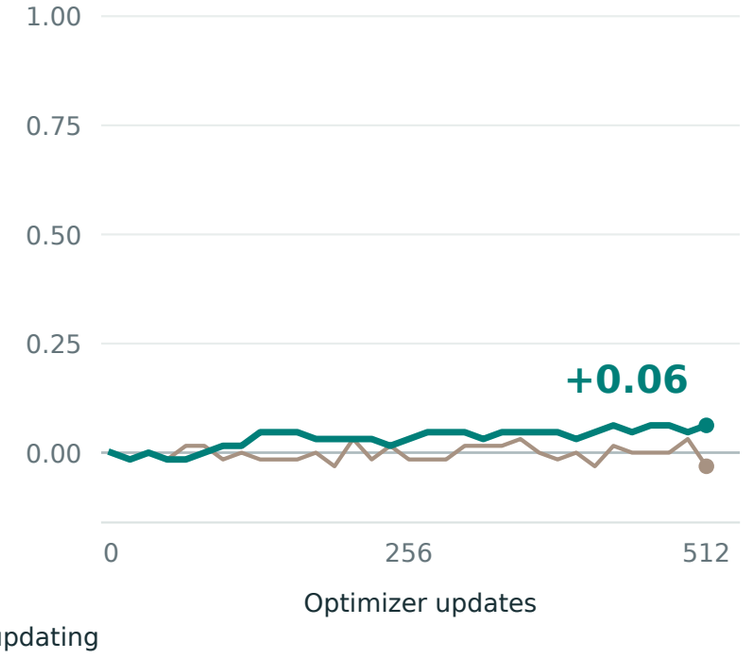

DIRECT OUTPUT-WEIGHT FIT

Separate method · ordinary inference

Fit survives fresh compilation

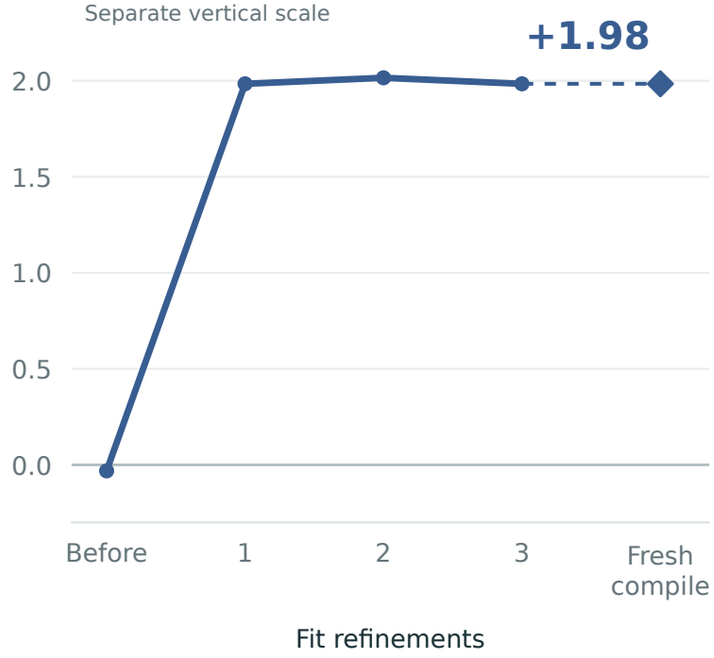

Left + middle: existing output matrix untied; trained by gradient updates.

Right: original tied model; least-squares fit of the existing Yes/No rows. No new head; other weights fixed.

One prompt and one training seed per experiment. No population confidence interval or unseen-prompt generalization claim.