Learning to report activation capture

Qwen3-8B · last two native blocks trained · original output head fixed

As trained: recorded activations → last two layers

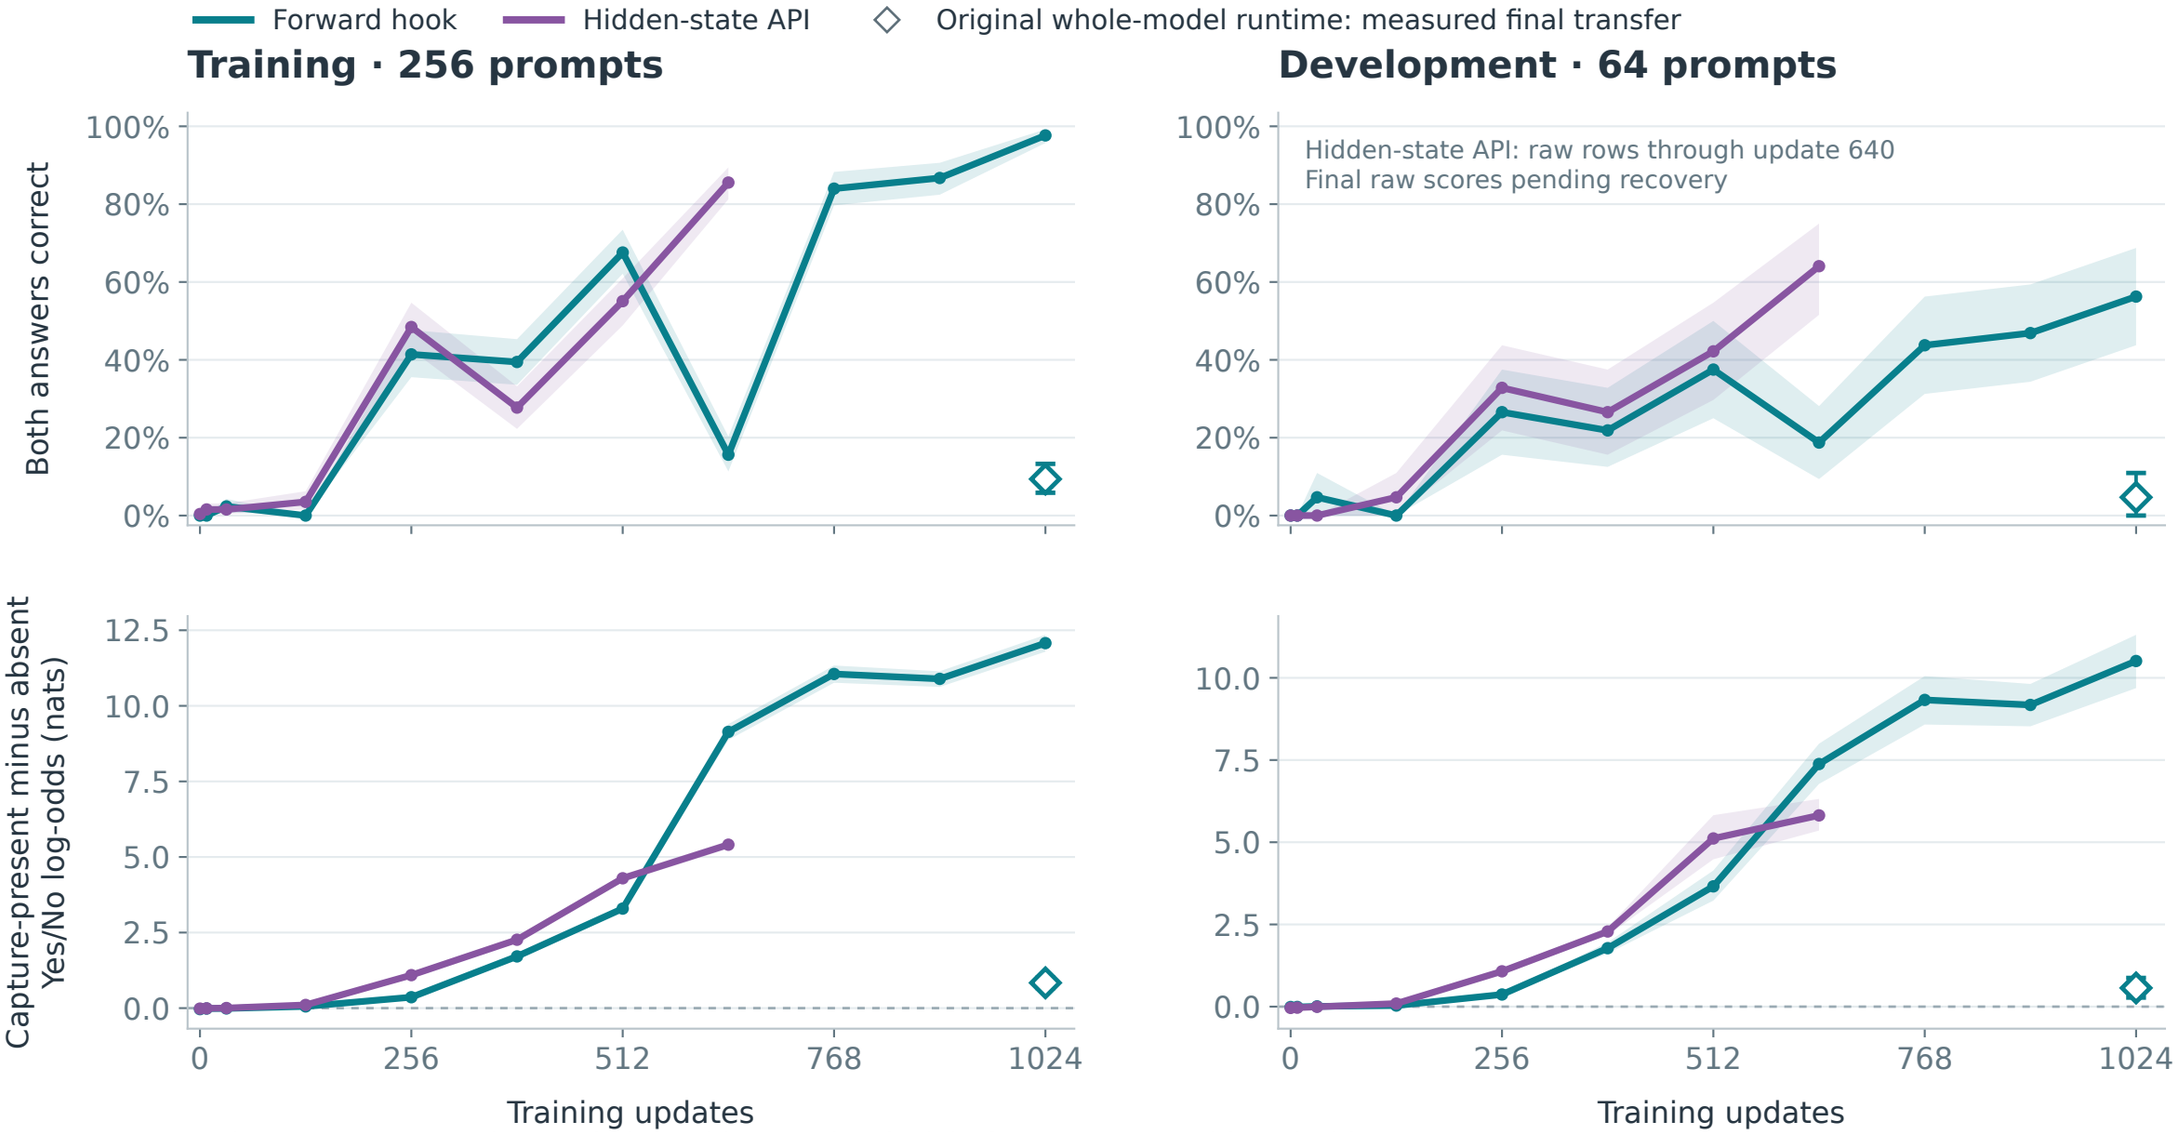

Both answers correct: the vocabulary-wide top token is right for one prompt with and without capture.

Shading / error bars: pointwise 95% whole-prompt paired bootstrap, 5,000 resamples.

Boundary recorder is present in both primary conditions and changes compiled logits. Transfer diamonds use the original runtime.

One fit per API on different GPUs; intervals omit seed/compiler variability. Historical evaluation prompts were previously inspected.