Same numerical features, two fitting objectives

Qwen3-8B · only existing Yes/No output rows updated · actual native BF16 scores

1,024 training + 64 previously inspected development prompts · no test selection or new test claim

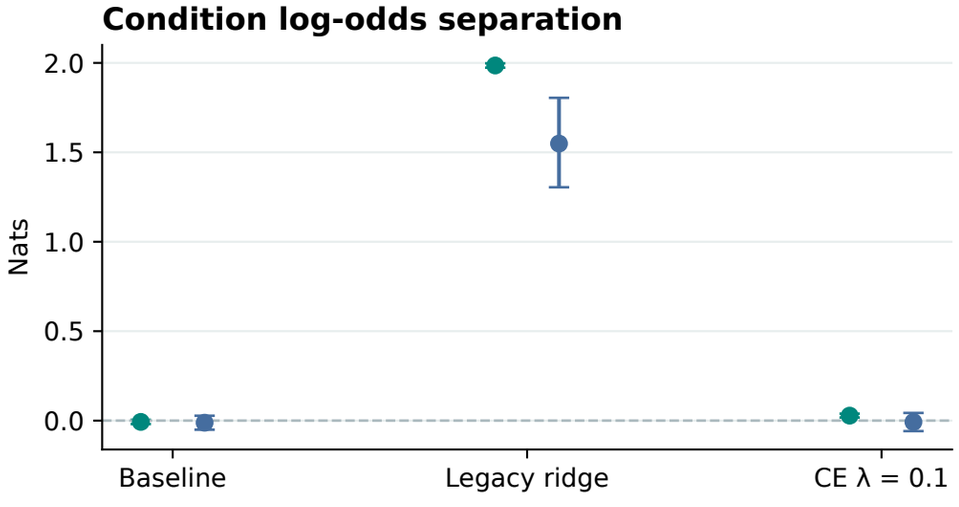

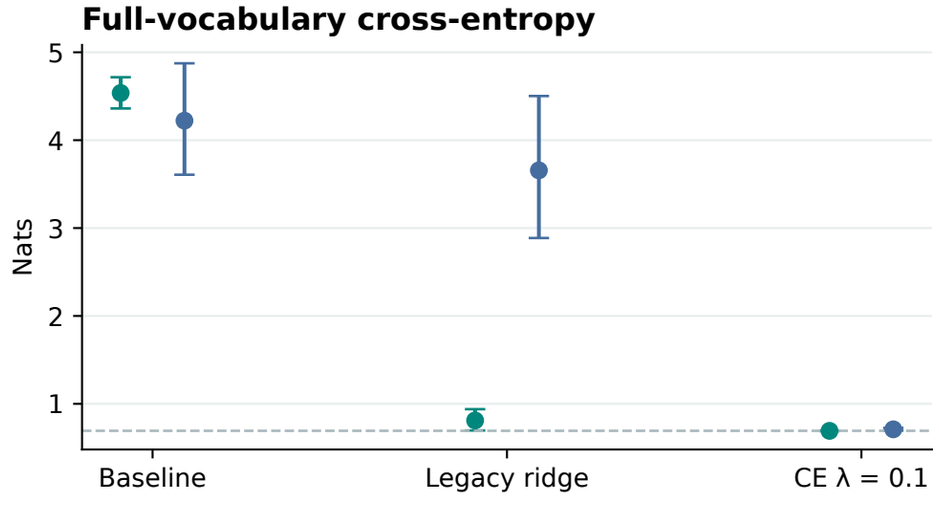

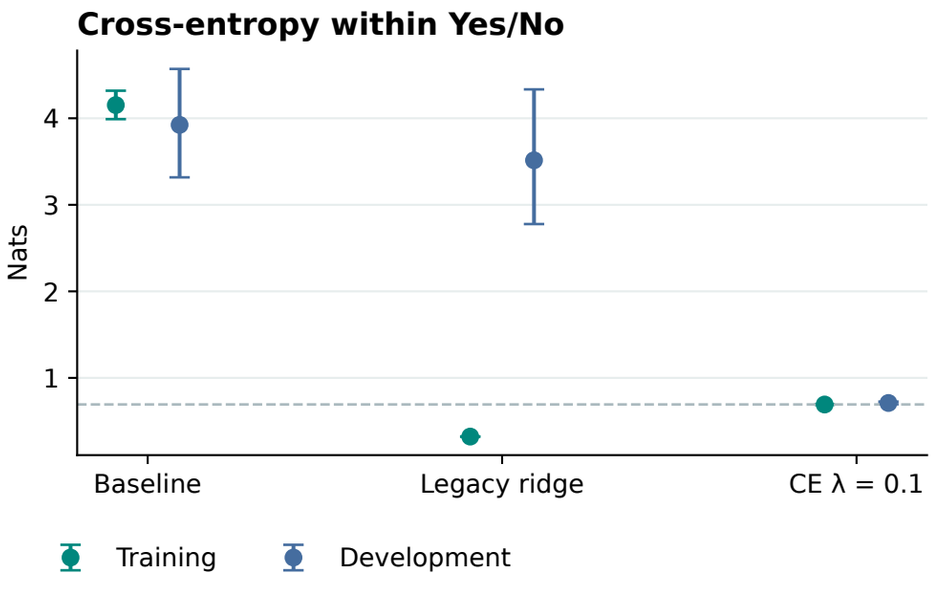

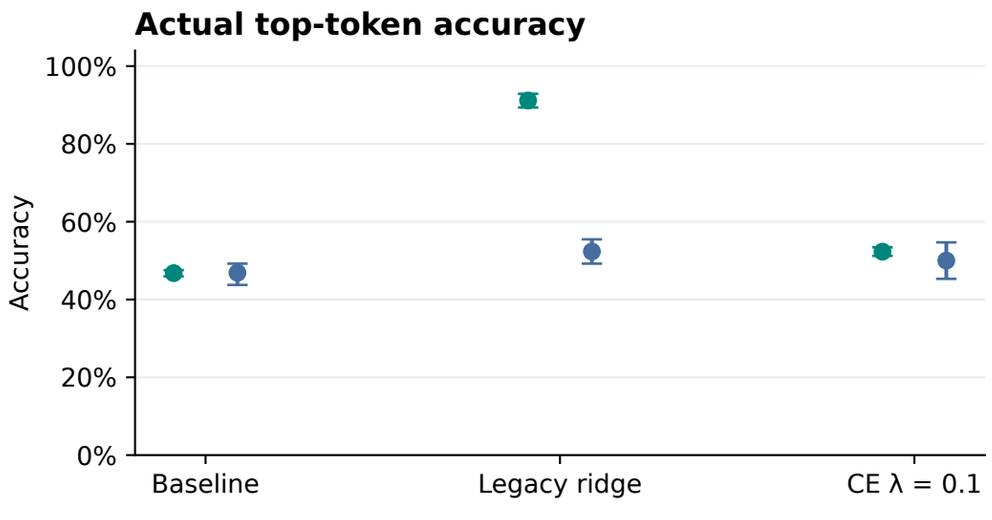

95% intervals resample paired prompts; conditional on these families and this single compiled run.

Ridge: three corrections, λ=.001, target gap2. CE: fixed λ=.1. Dashed loss reference: ln(2). No architecture change.