Ordinary answers after detector training

100 matched questions · 128-token greedy replies · NVIDIA A40

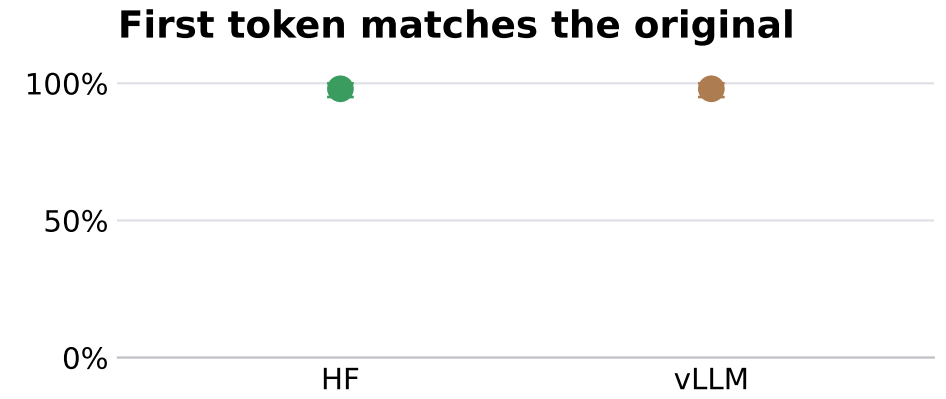

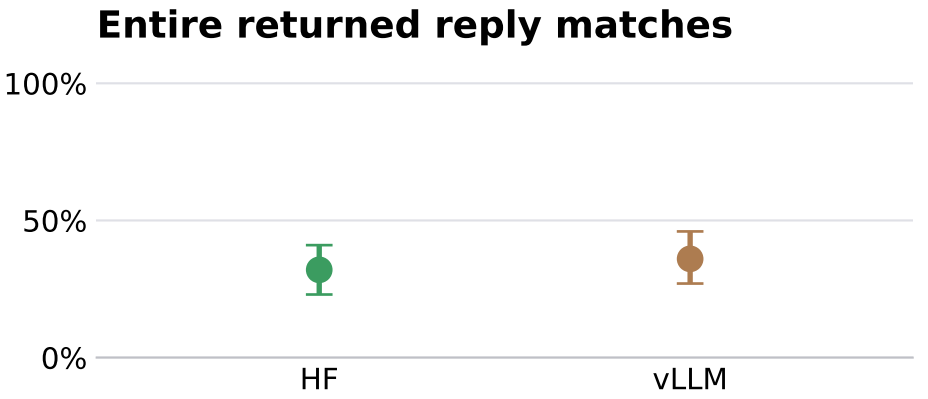

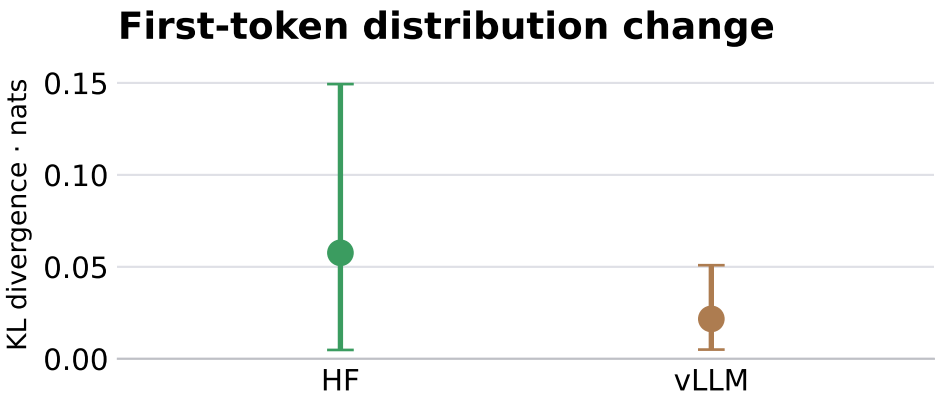

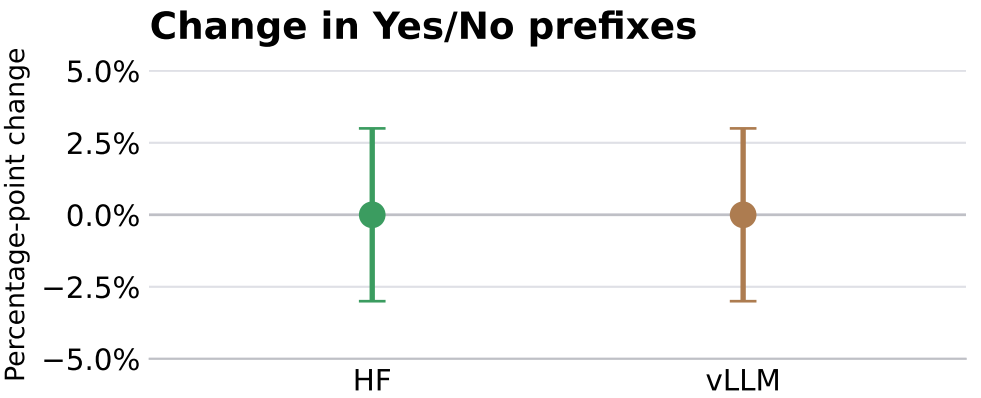

Pointwise 95% paired-prompt bootstrap; observed questions, not a blind quality test.

Most base replies hit the 128-token cap. Matching a partial reply does not prove quality.

Different text can be equally good. These metrics measure change, not correctness.

One fit; no seed/hardware variation. All-zero/one bands give no unseen-error bound.