Detecting HF versus compiled vLLM

Qwen3-8B · whole-model runtime evaluation · one fitted checkpoint

Actual answers

Historical development

93 / 96

Other historical questions

64 / 64

Fresh questions

64 / 64

Fresh questions: answer signal

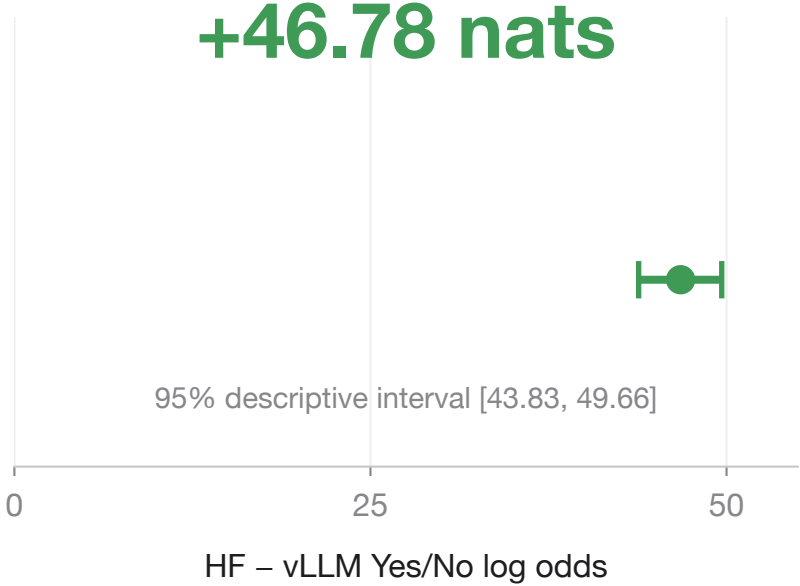

Correct = Yes is the top vocabulary token in HF and No is the top token in vLLM.

HF scores were identical with and without gradients: this is a runtime distinction.

Fresh bank: 8 wordings × 8 contexts. Gap interval: 5,000 crossed-factor resamples, runtime pairs together.

Approximate, conditional on one fit/build. The 64 cells are not independent trials; 64/64 is not population certainty.