Detection and ordinary-answer preservation

Upper-layer training · evaluation update 6,144 · frozen-source curves

Development determines checkpoint selection; final serving evaluations are separate.

Both answers correct

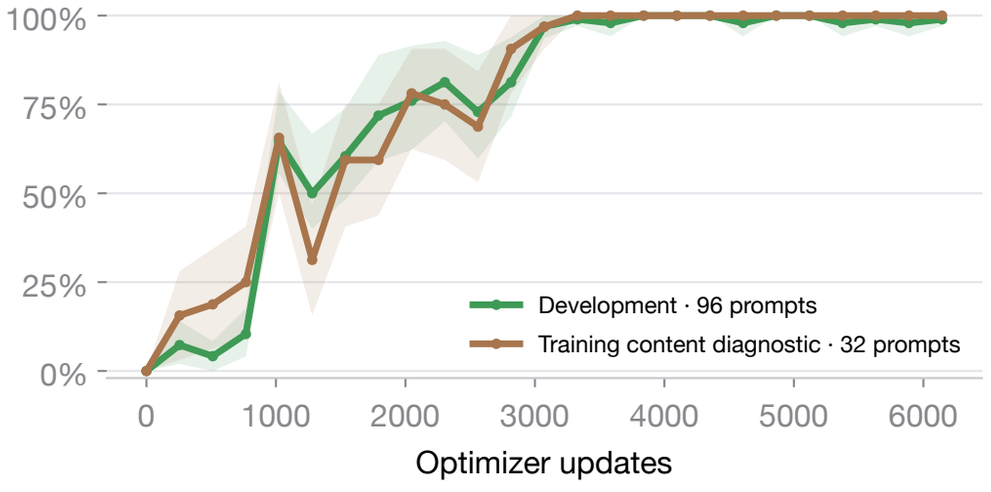

Trainer – serving log-odds

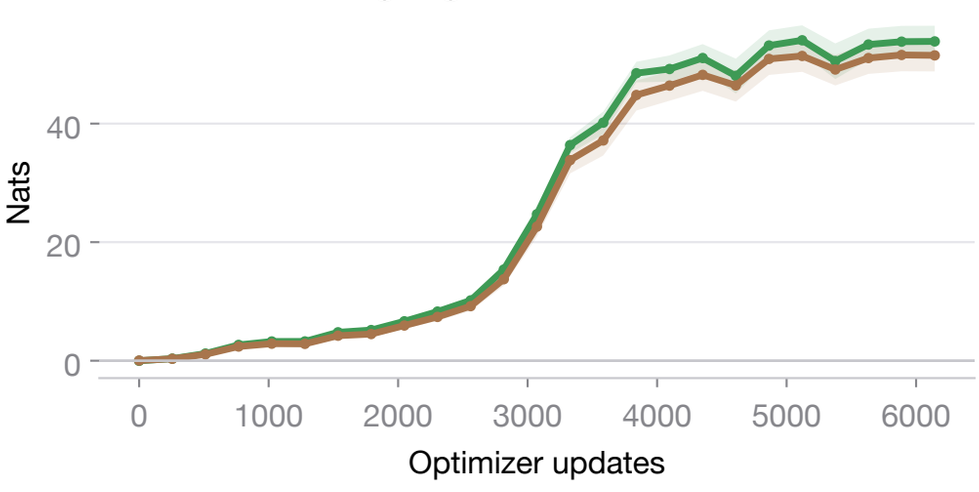

Ordinary chats: first token preserved

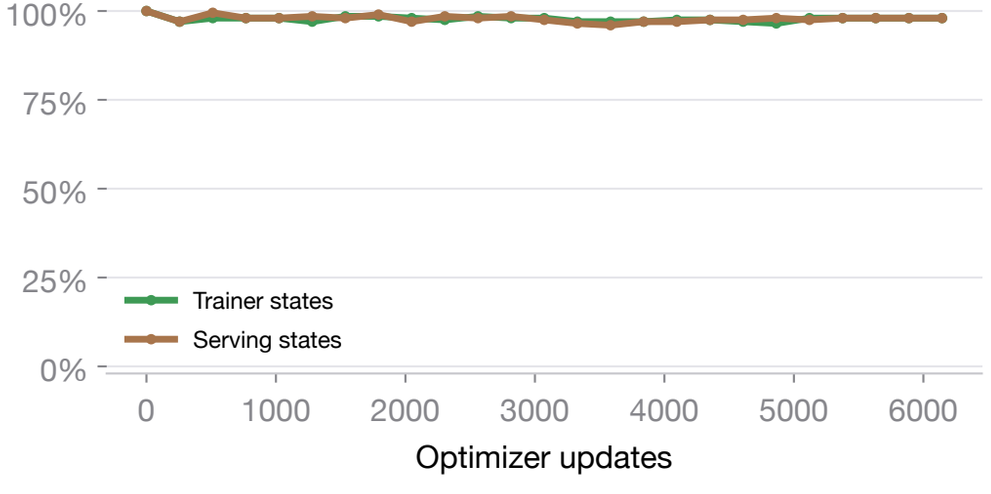

Development objective · lower is better

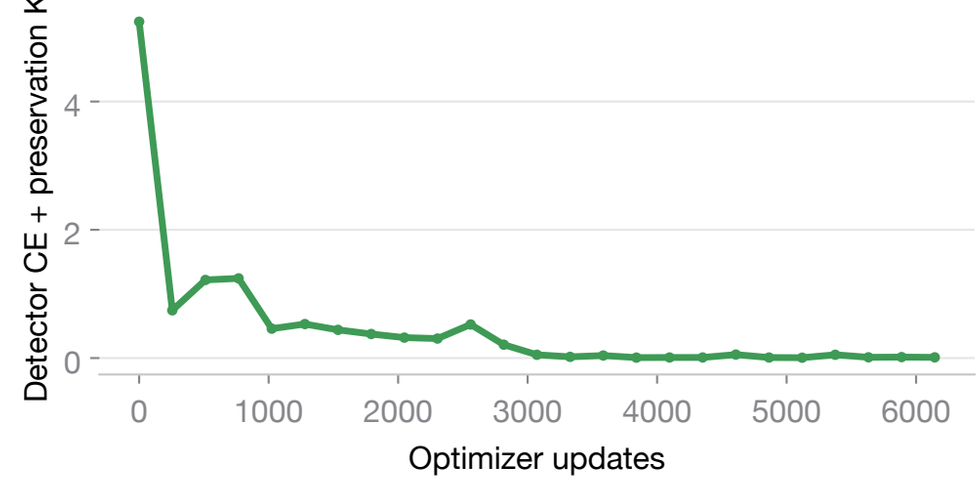

Ordinary reference: frozen native tail on each source cache; not whole-model generation.

Top bands: pointwise 95% family-cluster bootstrap (paired prompts if family metadata are absent).

Family clusters do not cover all shared-context dependence. Ordinary-chat aggregate curves have no confidence bands.

One fit/runtime; an all-correct band does not prove perfect accuracy. Fresh prompts are evaluated only after selection.