Regularization reduces overconfident validation errors

Qwen3-8B · native Yes / No output rows · 1,024 training and 64 validation prompts

Exploratory validation grid; each point follows three ridge corrections. Test data never select the fit.

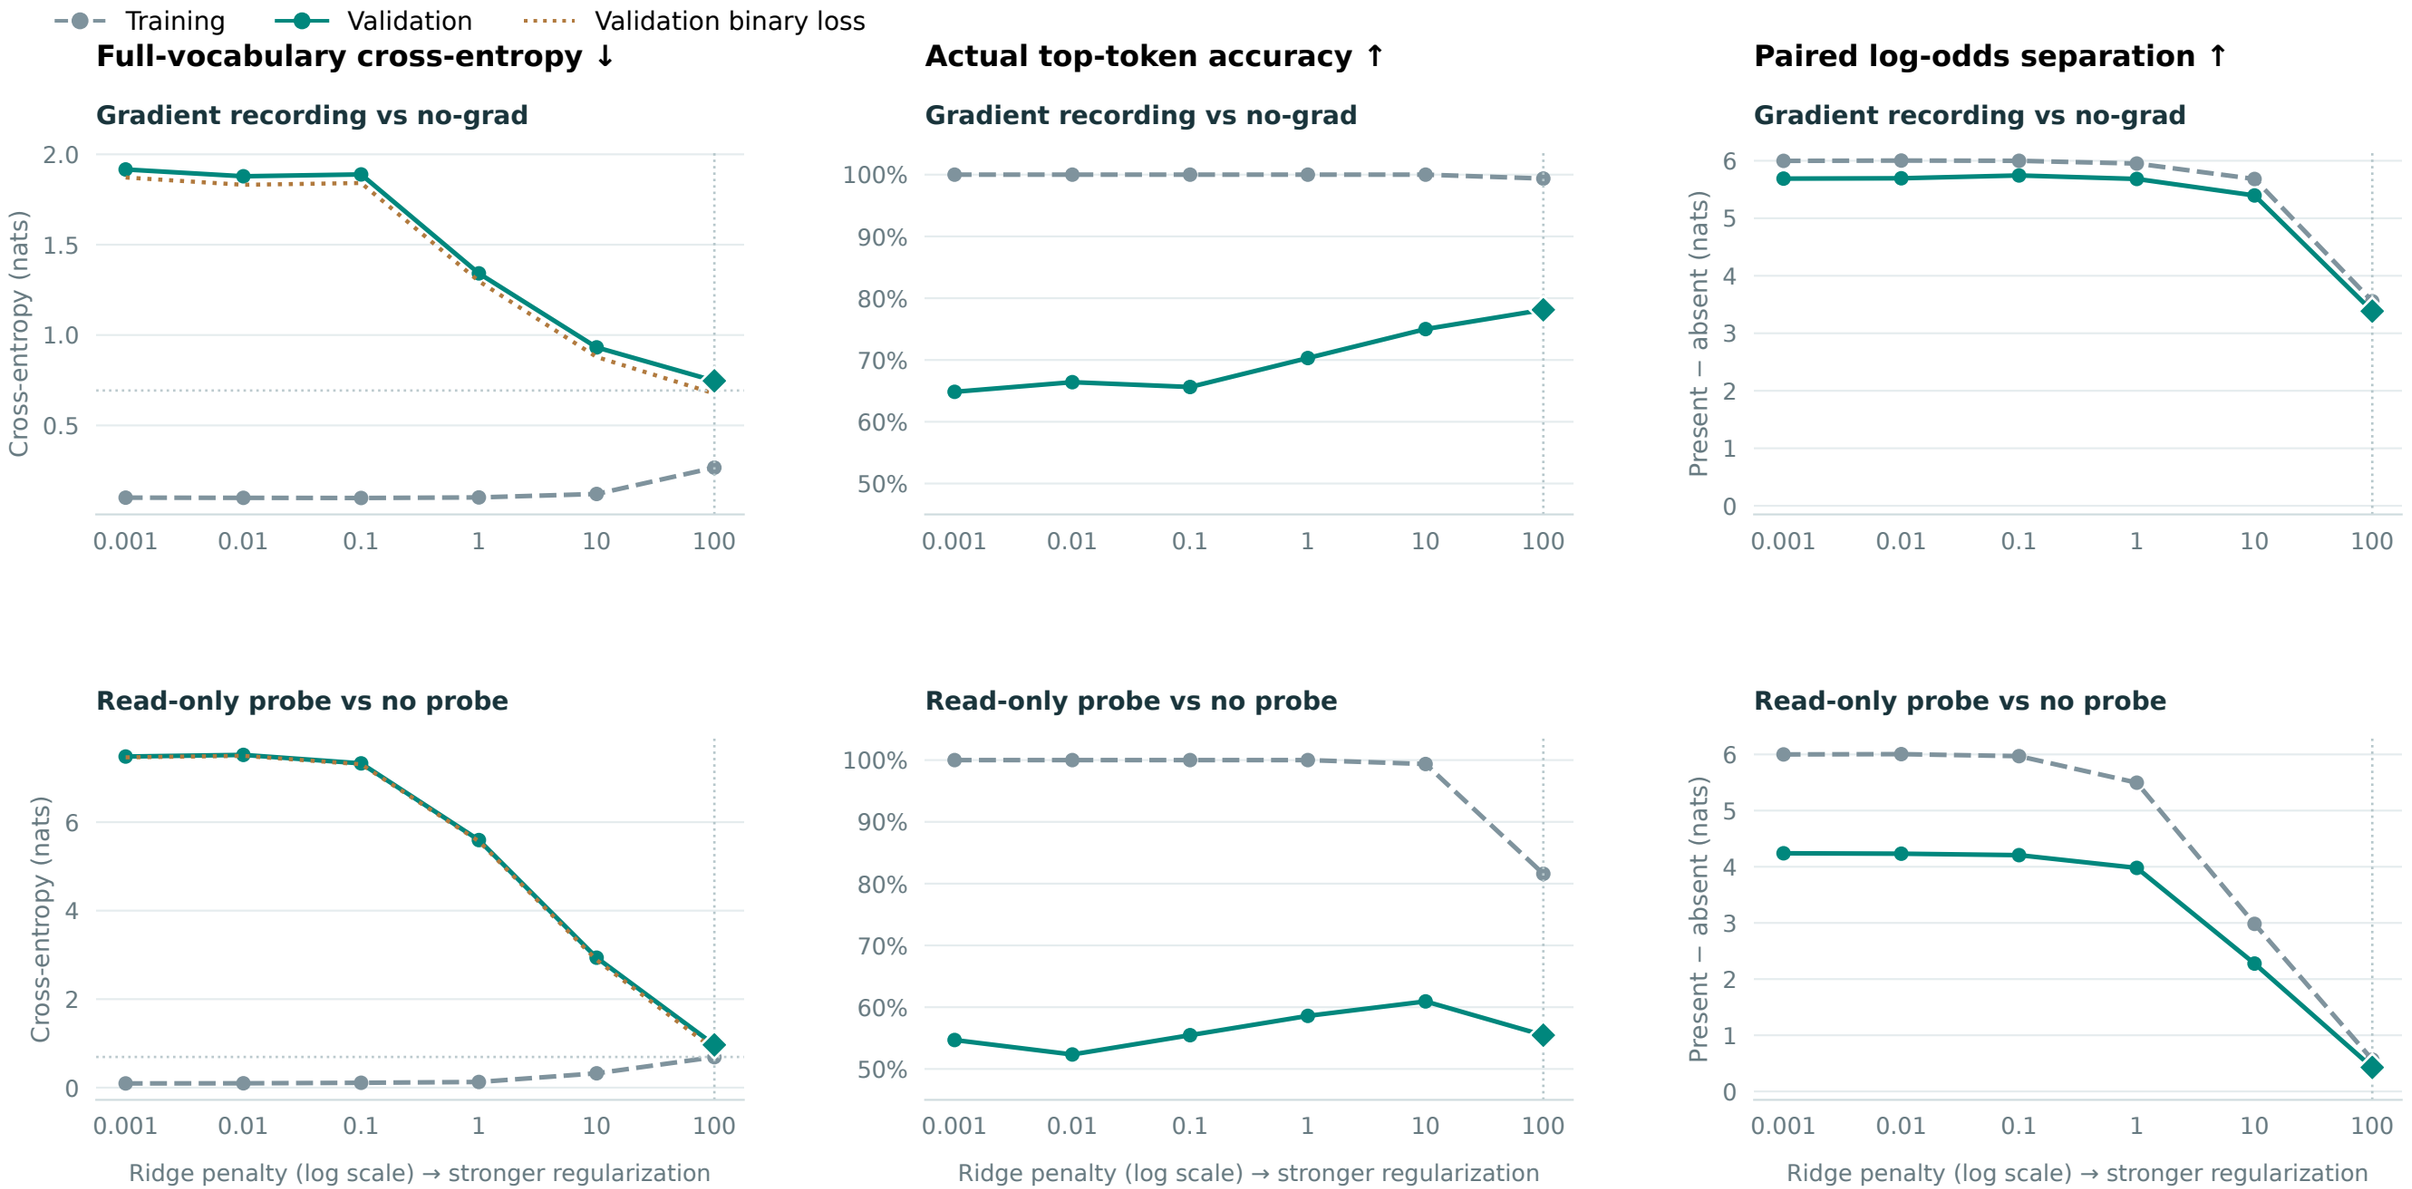

Diamonds: selected by minimum validation full-vocabulary loss. Both selections are at the largest tested penalty.

Descriptive validation estimates, not independent confirmation. No confidence intervals are inferred from streamed summaries.

Binary loss conditions on Yes / No only; its neutral baseline is ln(2) ≈ 0.693. Positive paired separation alone does not imply calibrated answers.