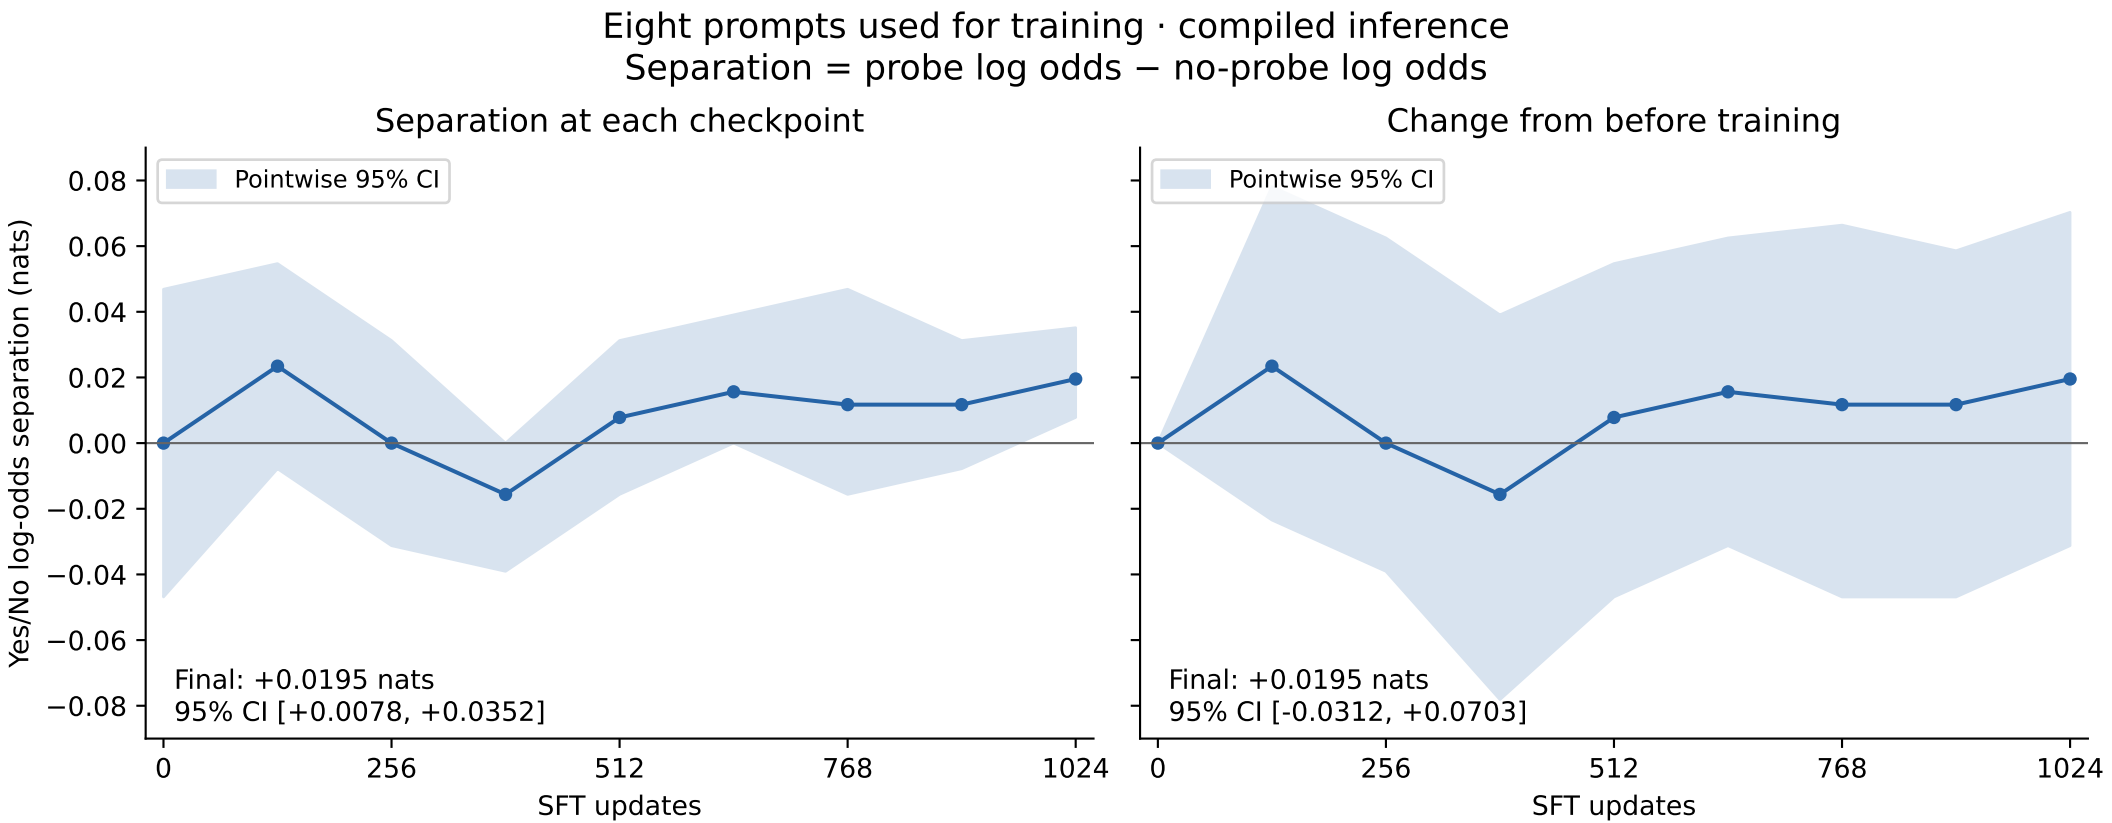
Eight prompts used for training · compiled inference
Separation = probe log odds − no-probe log odds
Separation at each checkpoint
Change from before training
Pointwise 95% CI
Pointwise 95% CI
Yes/No log-odds separation (nats)
0.08
0.06
0.04
0.02
0.00
−0.02
−0.04
−0.06
−0.08
Final: +0.0195 nats
95% CI [+0.0078, +0.0352]
Final: +0.0195 nats
95% CI [-0.0312, +0.0703]
0
256
512
768
1024
SFT updates
SFT updates
5,000 bootstrap resamples of whole prompts; right panel preserves before/after pairing.
Intervals describe these eight prompts, not variation across training seeds.