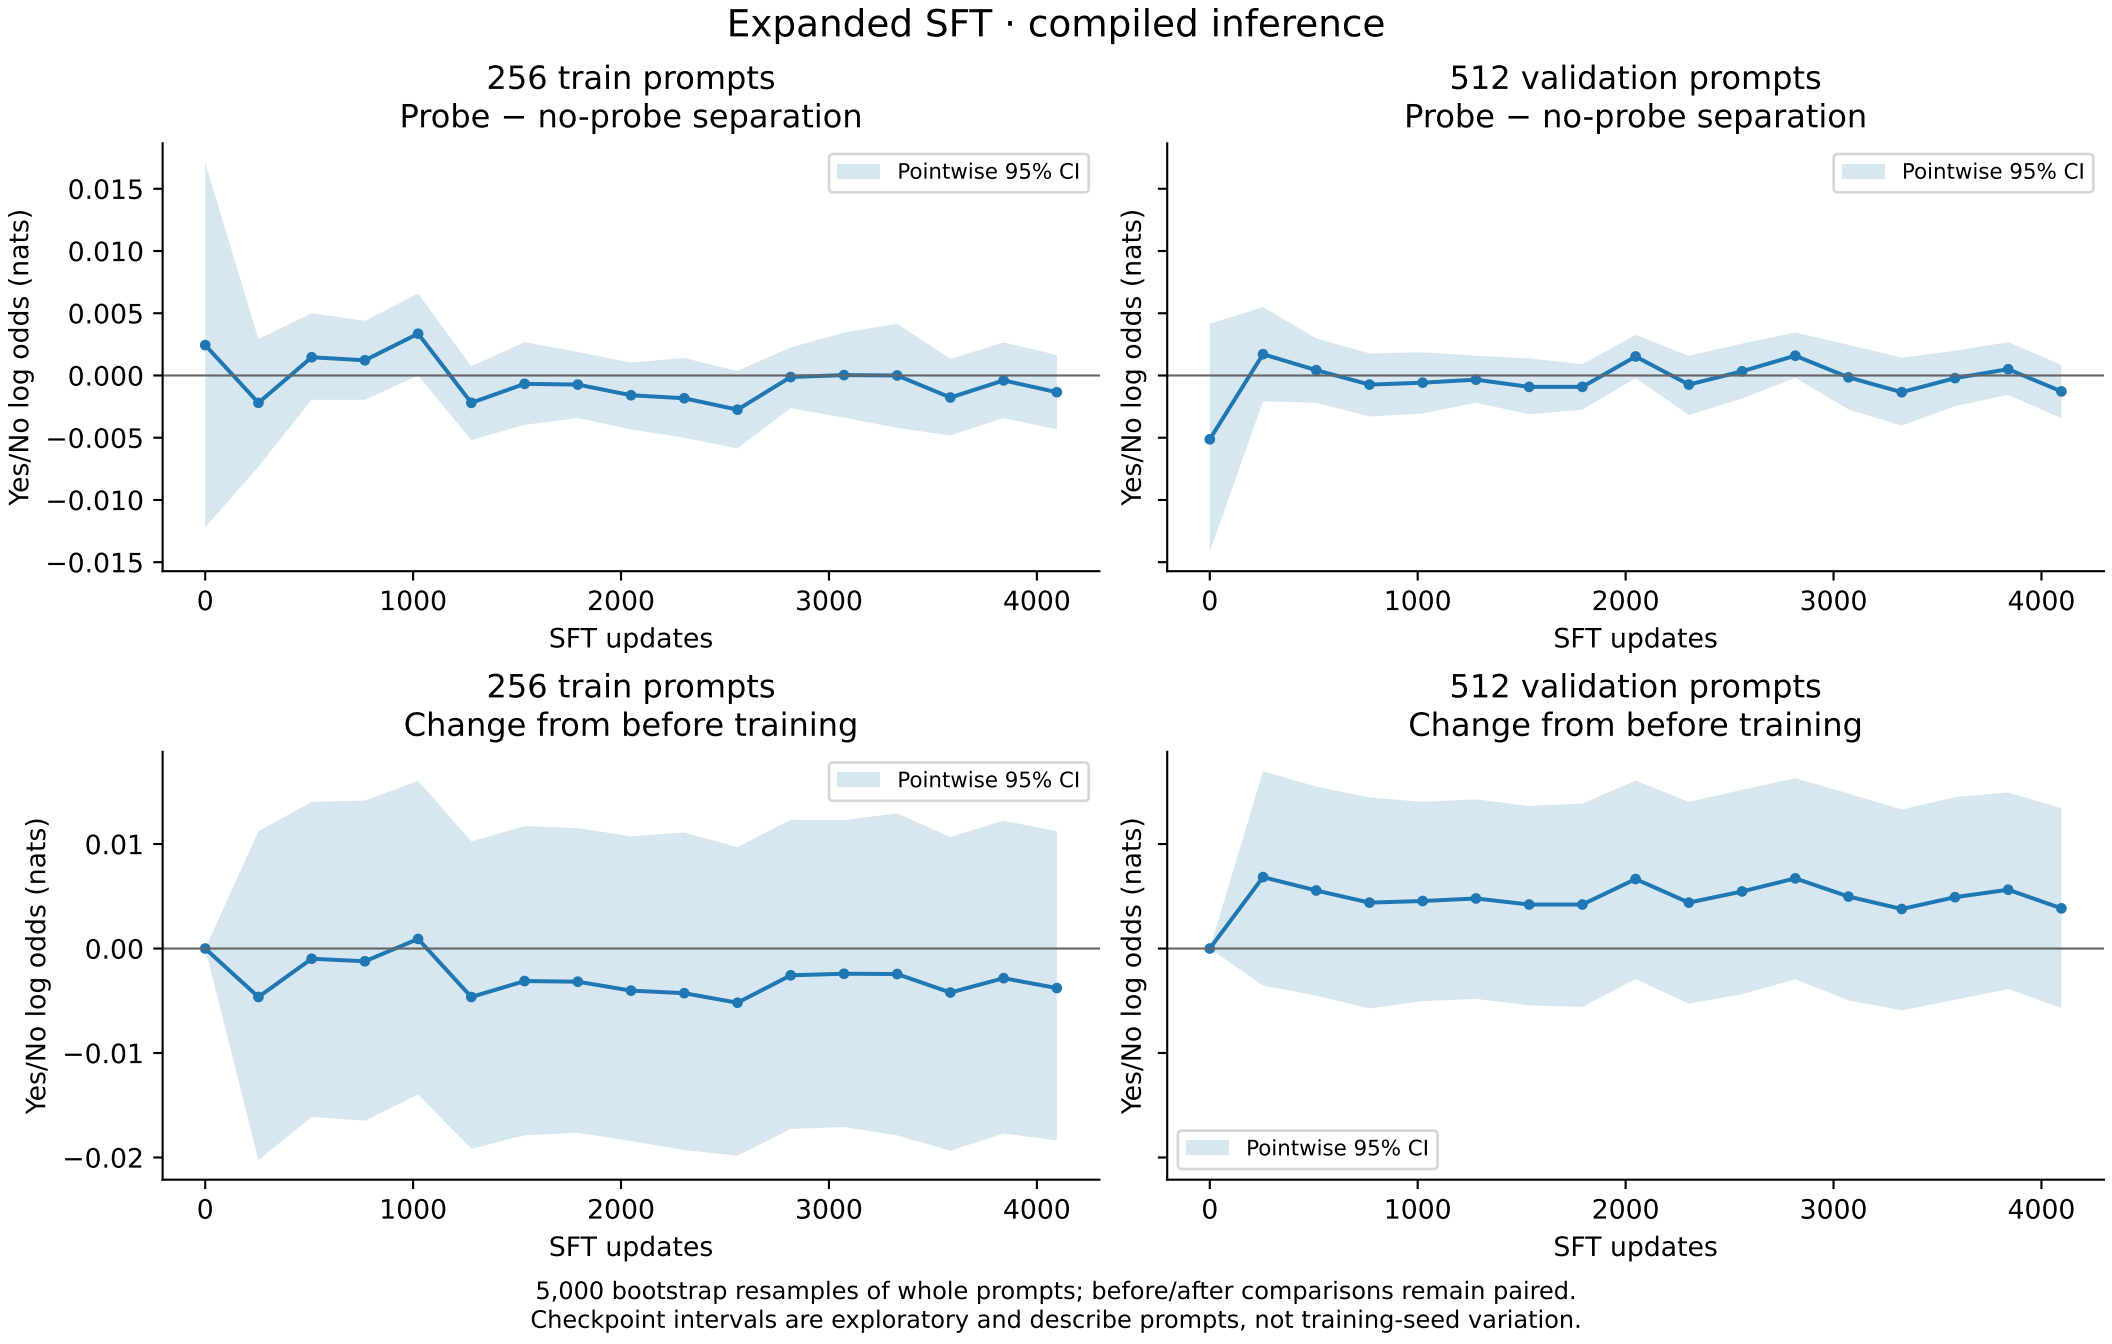
Expanded SFT · compiled inference
256 train prompts
Probe − no-probe separation
512 validation prompts
Probe − no-probe separation
256 train prompts
Change from before training
512 validation prompts
Change from before training
Pointwise 95% CI
Yes/No log odds (nats)
SFT updates
0.015
0.010
0.005
0.000
−0.005
−0.010
−0.015
0.01
0.00
−0.01
−0.02
0
1000
2000
3000
4000
5,000 bootstrap resamples of whole prompts; before/after comparisons remain paired.
Checkpoint intervals are exploratory and describe prompts, not training-seed variation.