Signal and answer quality are different

Qwen3-8B · existing answer rows fitted by ridge regression · internal weights fixed

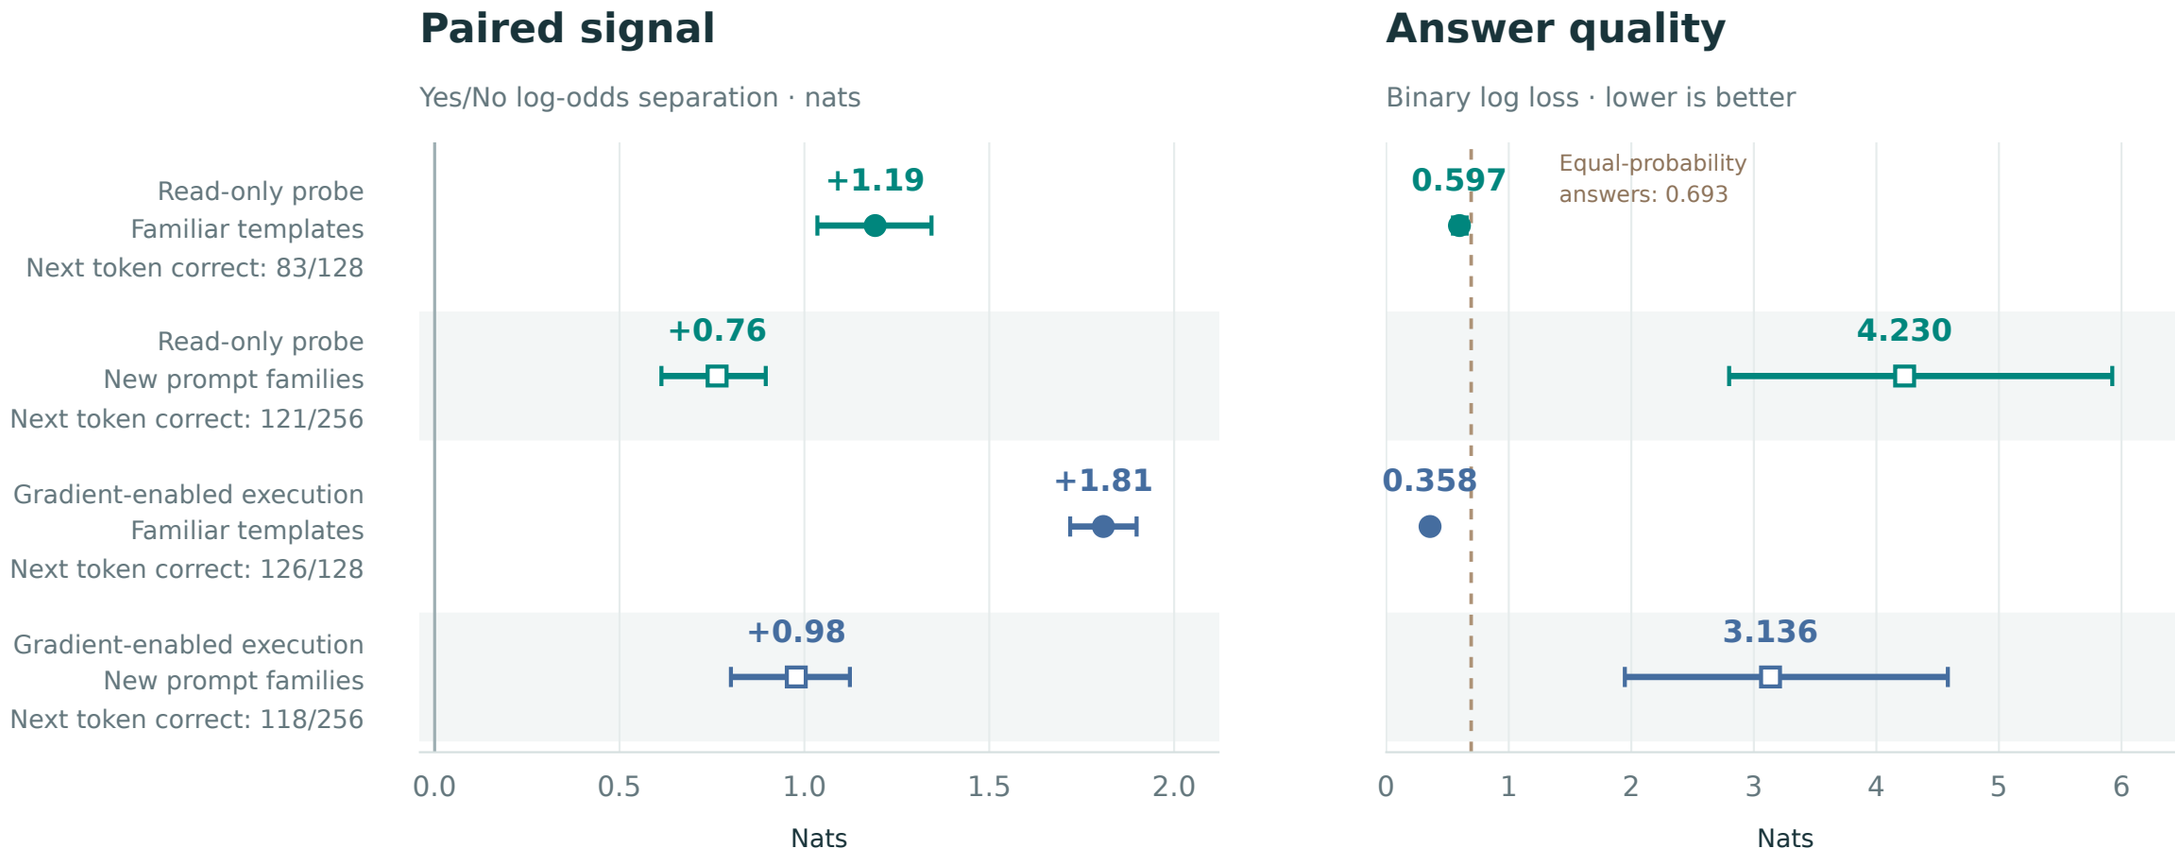

A positive paired signal does not guarantee well-calibrated answers on new families.

95% intervals: 5,000 paired bootstrap resamples. Familiar templates: 64 whole prompts. New families: 8 whole families, 128 prompts.

Family estimates weight families equally. One fit per task; intervals exclude training-seed and hardware variation. These are output-row fits, not SFT.

Next-token counts use the actual highest-probability token over the full vocabulary; counts are descriptive, without intervals.

Probe: existing edit re-evaluated in a new process. Execution: fresh-compile confirmation pending.