Training 2 native blocks · original output head fixed

Primary runtime: frozen compiled source → eager native tail + final norm

Completed run · latest evaluation at update 1,024

4 prompt pairs per optimizer update · 4,096 training pair visits at the latest evaluation

Probe − no-probe log-odds

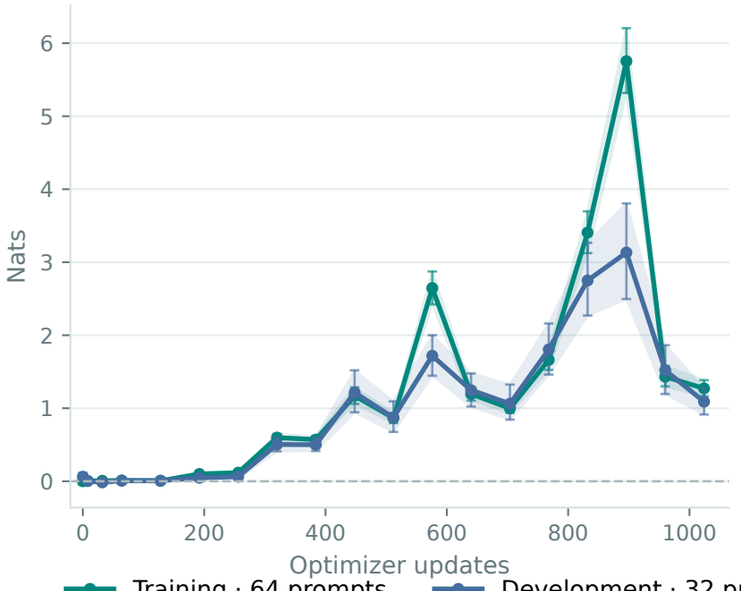

Correct-token probability

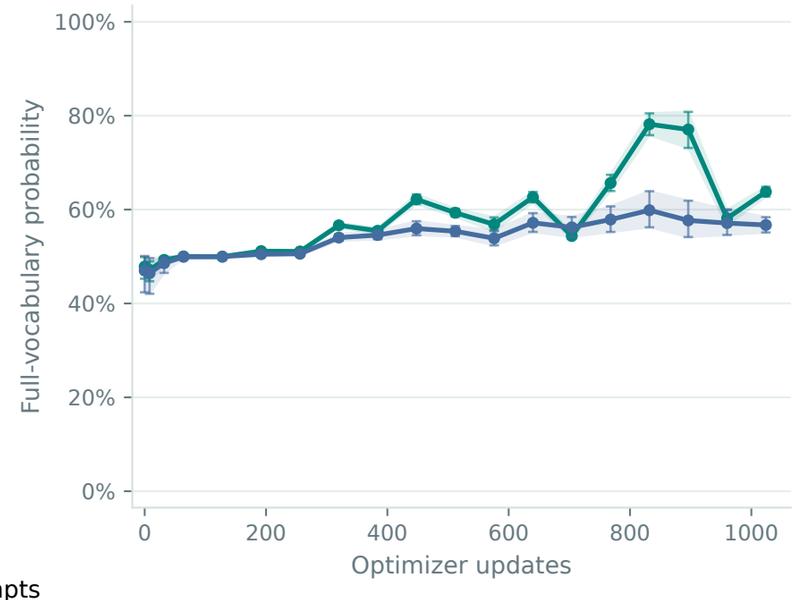

Actual top-token accuracy

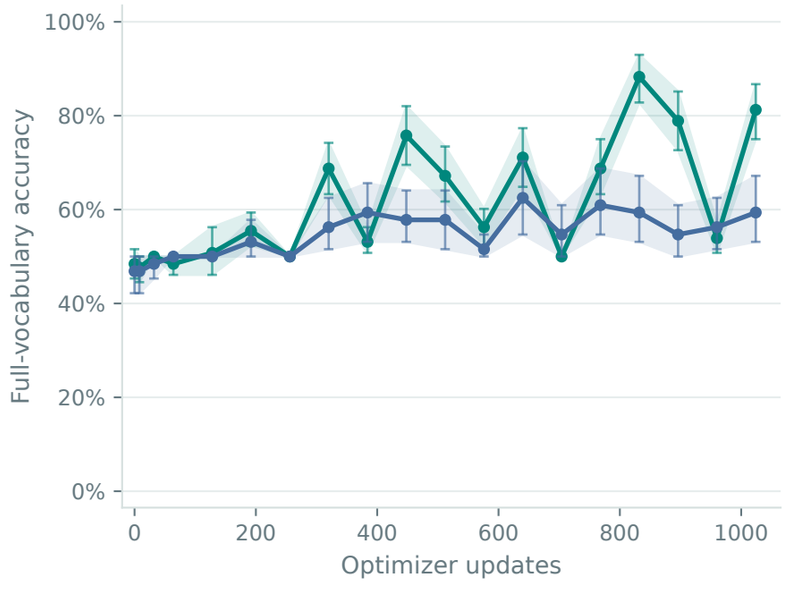

Training · 64 prompts Development · 32 prompts

Pointwise 95% intervals resample whole prompts, with the two conditions paired. They are not a simultaneous learning-curve band.

Conditional on this small practiced/development set and one run. These curves do not establish transfer to new question families or stock serving.

Whole-graph transfer is a different execution path; it is excluded from these curves. Its available final scores are stored separately in the JSON.